

Extending Algebraic Multigrid Solvers

SciDAC 11

Luke Olson, University of Illinois at Urbana-Champaign

Jacob Schroder	Colorado (LLNL, SNL)	generalized interpolation
Hormozd Gahvari	Illinois (LLNL)	multigrid scalability
James Lai	Illinois (SNL)	high-order conforming
Steven Dalton	Illinois (LLNL, Nvidia)	accelerating multigrid
John Hall	Illinois (LLNL)	automating multigrid
Bill Gropp	Illinois	
Rob Falgout	LLNL	
Ulrike Meier-Yang	LLNL	
Ray Tuminaro	SNL	
Scott MacLachlan	Tufts	

$Ax = b$: old, but persistent

- large, sparse matrix
- represents a major bottleneck in many simulations both small and large
- advanced Krylov solvers + fancy **preconditioning**

the obvious trend...

- complex, non-symmetric, unstructured, not-so-nice applications
- bio, information, high-dimension, stochastic, etc.
- heterogeneous, large-scale architectures
- thinner cores, acceleration units

the obvious trend...

- complex, non-symmetric, unstructured, not-so-nice applications
- bio, information, high-dimension, stochastic, etc.



elliptic not motivating

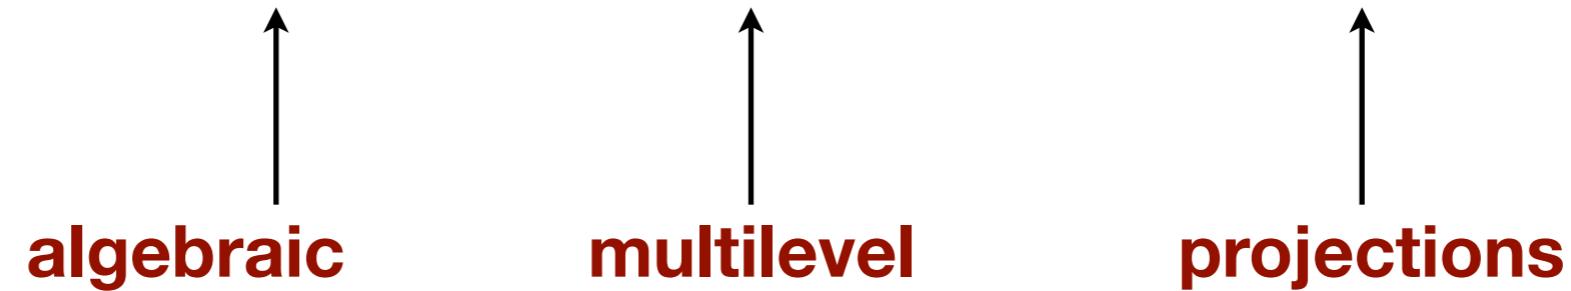
- heterogeneous, large-scale architectures
- thinner cores, acceleration units



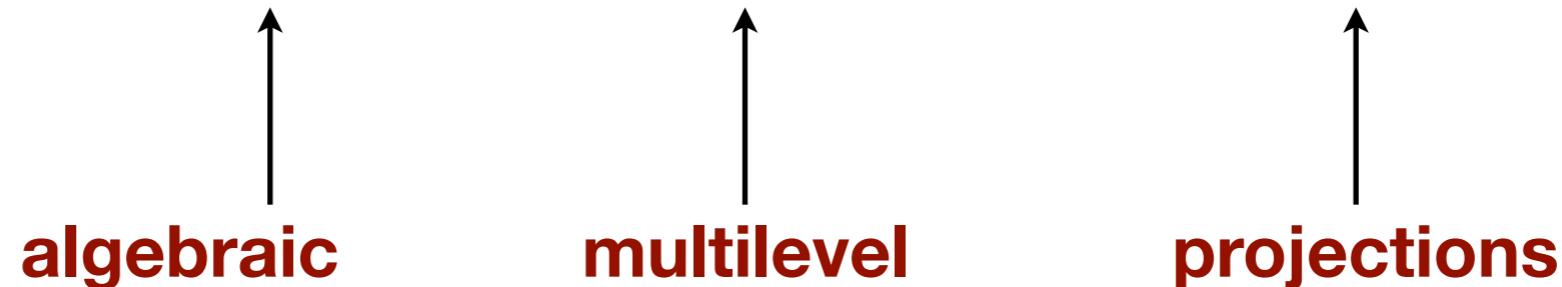
coarse and fine
grained parallelism

wanted: more robust, efficient, and principled solver

wanted: more robust, efficient, and principled solver



wanted: more robust, efficient, and principled solver



solution: move beyond basic problems

solution: consider communication/complexity

solution: optimality into the multilevel hierarchy

show some success in robustness

highlight optimality

progress in high-performance

algebraic multilevel methods are flexible

“...but multigrid is old”

no!

multigrid is only a structure; new
approaches are transforming
solver development

yes!

history of successful development
and continued progress

Multilevel view

attenuate high energy quickly with relaxation
attenuate low energy error through coarse-grid correction

Multilevel view

Jacobi, Gauss-Seidel
Chebychev
ILU
Kaczmarz

attenuate high energy quickly with relaxation
attenuate low energy error through coarse-grid correction

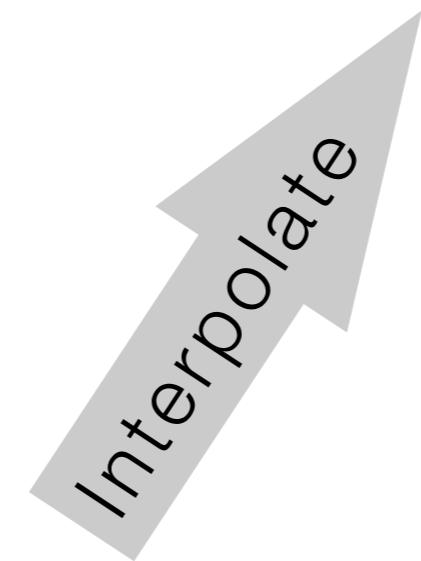
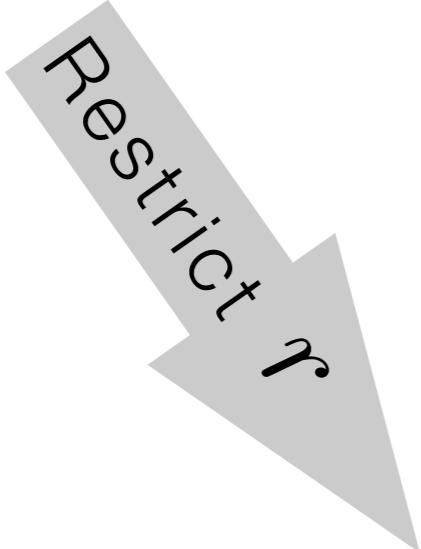
DD
projection methods
Interpolation

Multilevel view

attenuate high energy quickly with relaxation
attenuate low energy error through coarse-grid correction

relax $Ax = b$

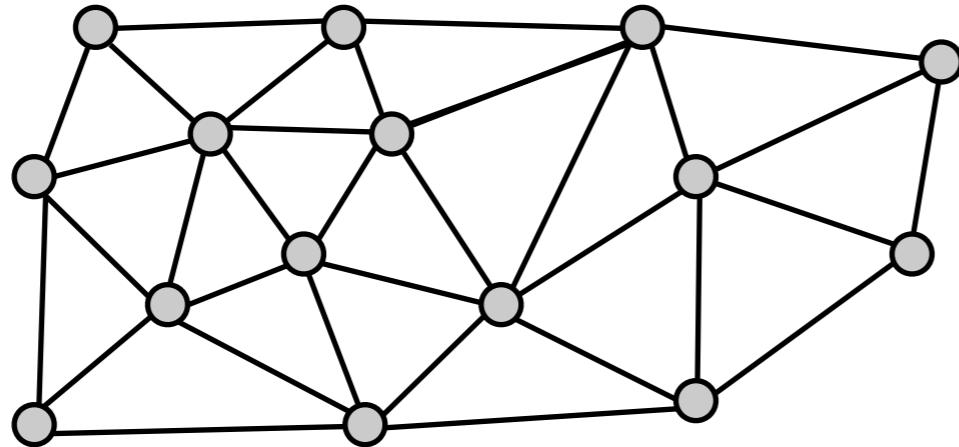
correct $x \leftarrow x + Pe_c$



solve coarse problem $A_c e_c = r_c$

Algebraic Framework

- aggregation: groups of fine nodes form coarse nodes

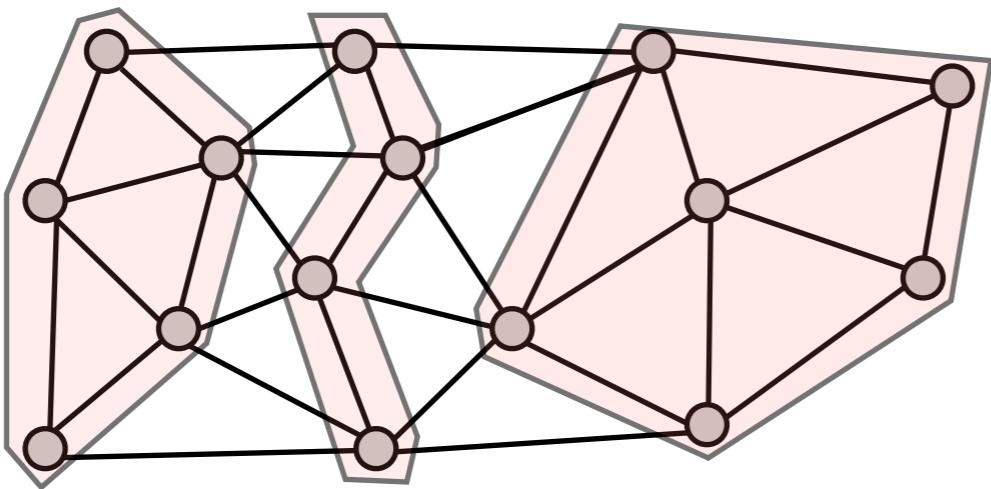


- an initial interpolation pattern
- find an optimal interpolation operator P that contains low energy
- cycle:
$$x \leftarrow x + P(P^T A P)^{-1} P^T r$$

$$x \leftarrow x + P A_{coarse}^{-1} P^T r$$

Algebraic Framework

- aggregation: groups of fine nodes form coarse nodes



fine: 15
coarse: 3

- an initial interpolation pattern
- find an optimal interpolation operator P that contains low energy

- cycle:
$$x \leftarrow x + P(P^T A P)^{-1} P^T r$$

$$x \leftarrow x + P A_{coarse}^{-1} P^T r$$

what we need

- an idea of the low energy: physics, training, intuition

$$AB \approx 0$$

- a strength measure to determine strong node couplings

$$S_{ij}$$

- a parallel aggregation method

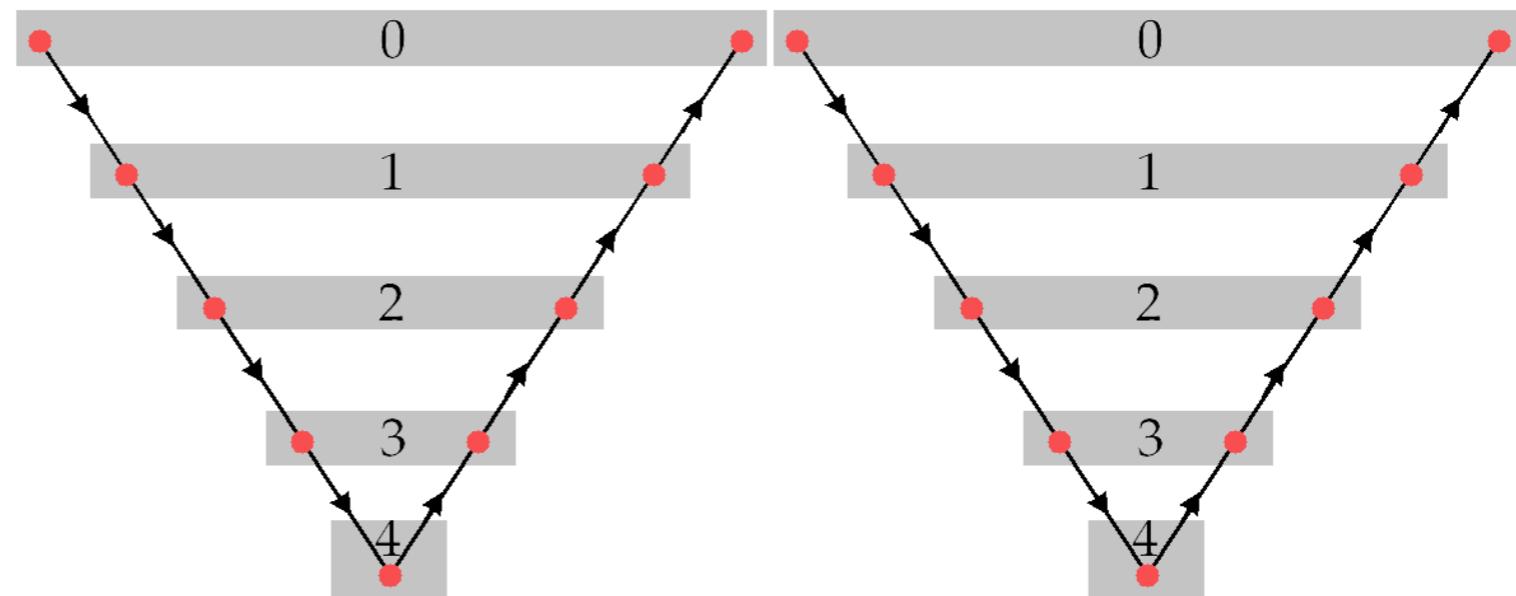
$$Agg$$

- low complexity, optimal interpolation

$$P$$

- better cycling

- richer coarse grids
- “parallel” cycling

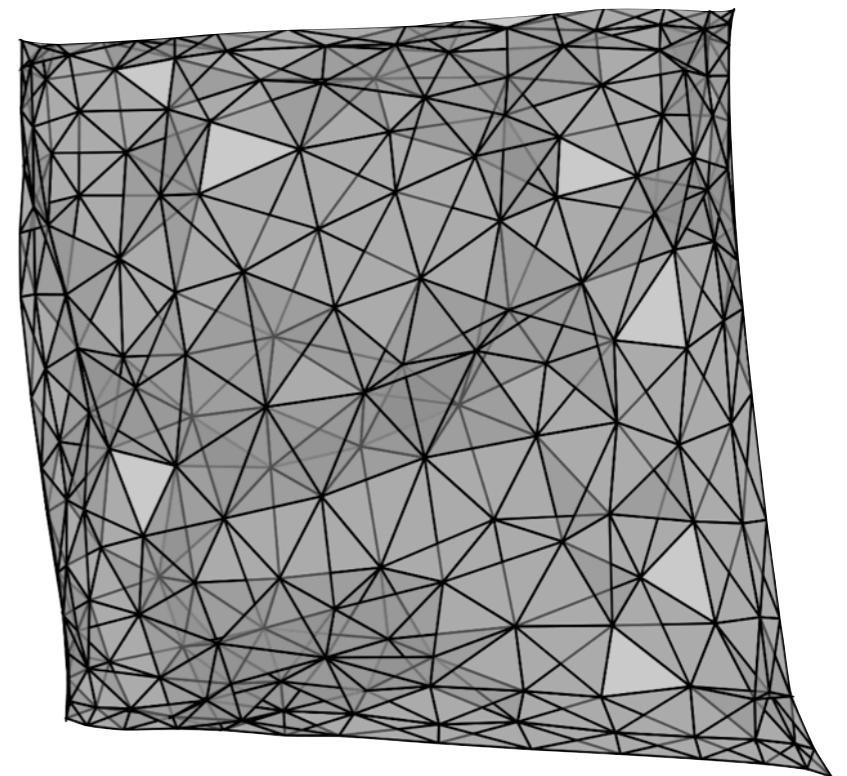
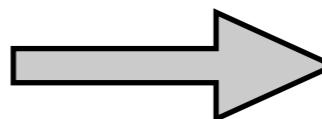
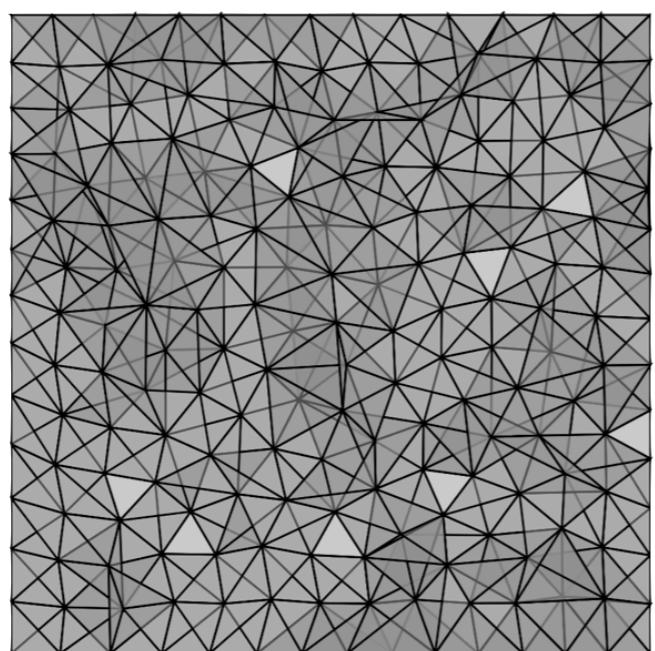
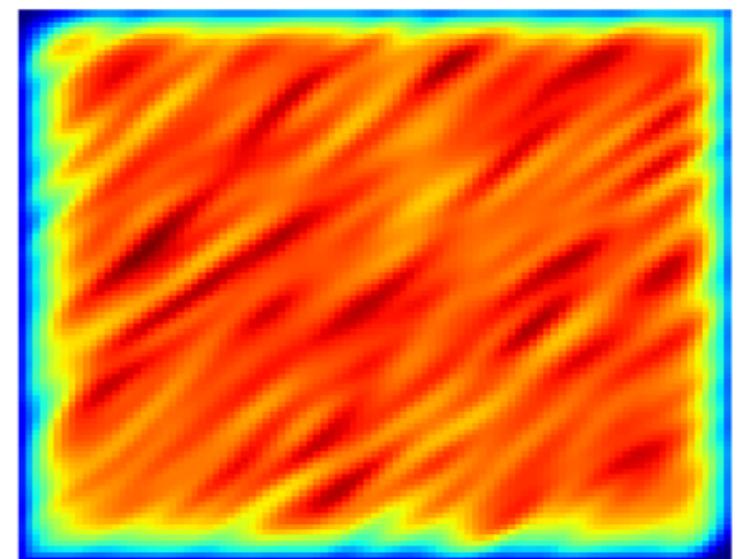
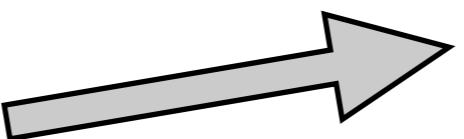


Setup

Solve

Obtaining Low-energy

- *a priori* knowledge
- adapting a cycle*
- pre-smooth guesses
- topological inference



Classic Strength: the first mistake

- i strongly depends on j if

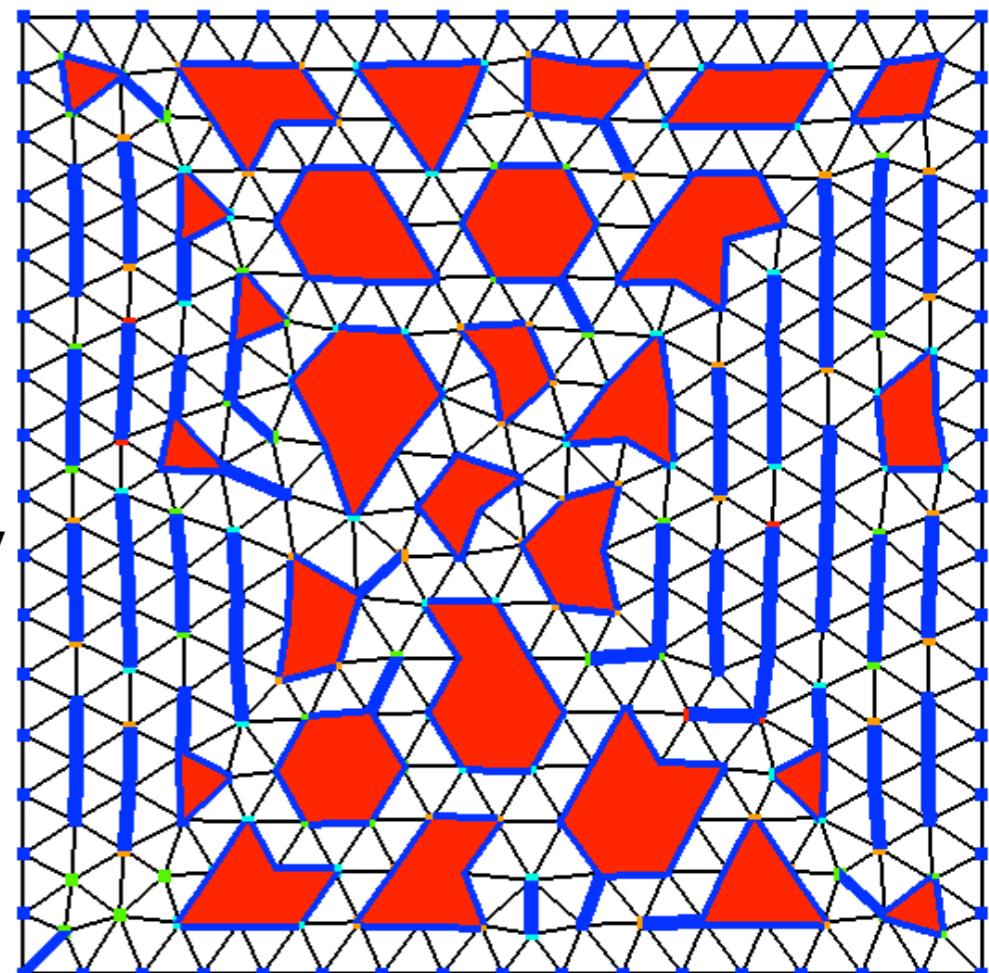
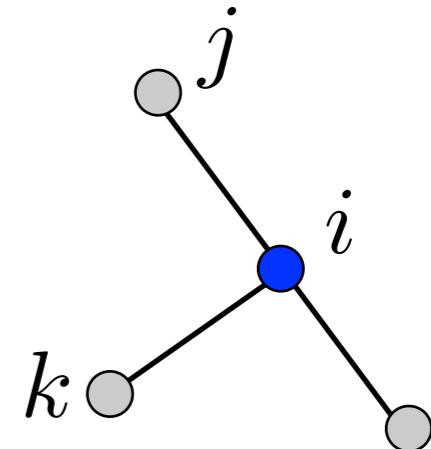
$$\frac{|A_{ij}|}{\sqrt{A_{ii}A_{jj}}} \geq tol$$

- i strongly depends on j if

$$-A_{ij} \geq tol * \max_{k \neq i} -A_{ik}$$

both “think” elliptic

↑
anisotropy
↓

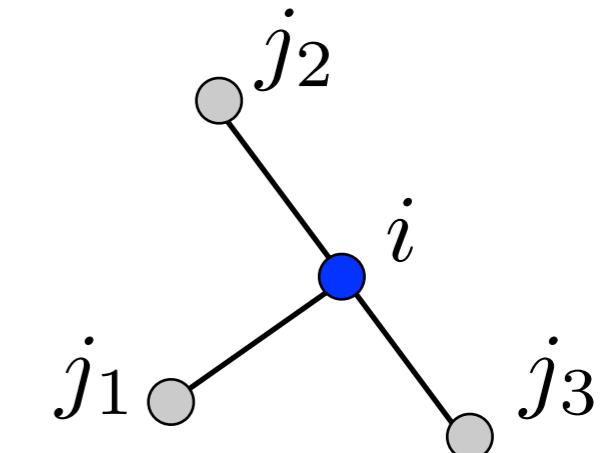
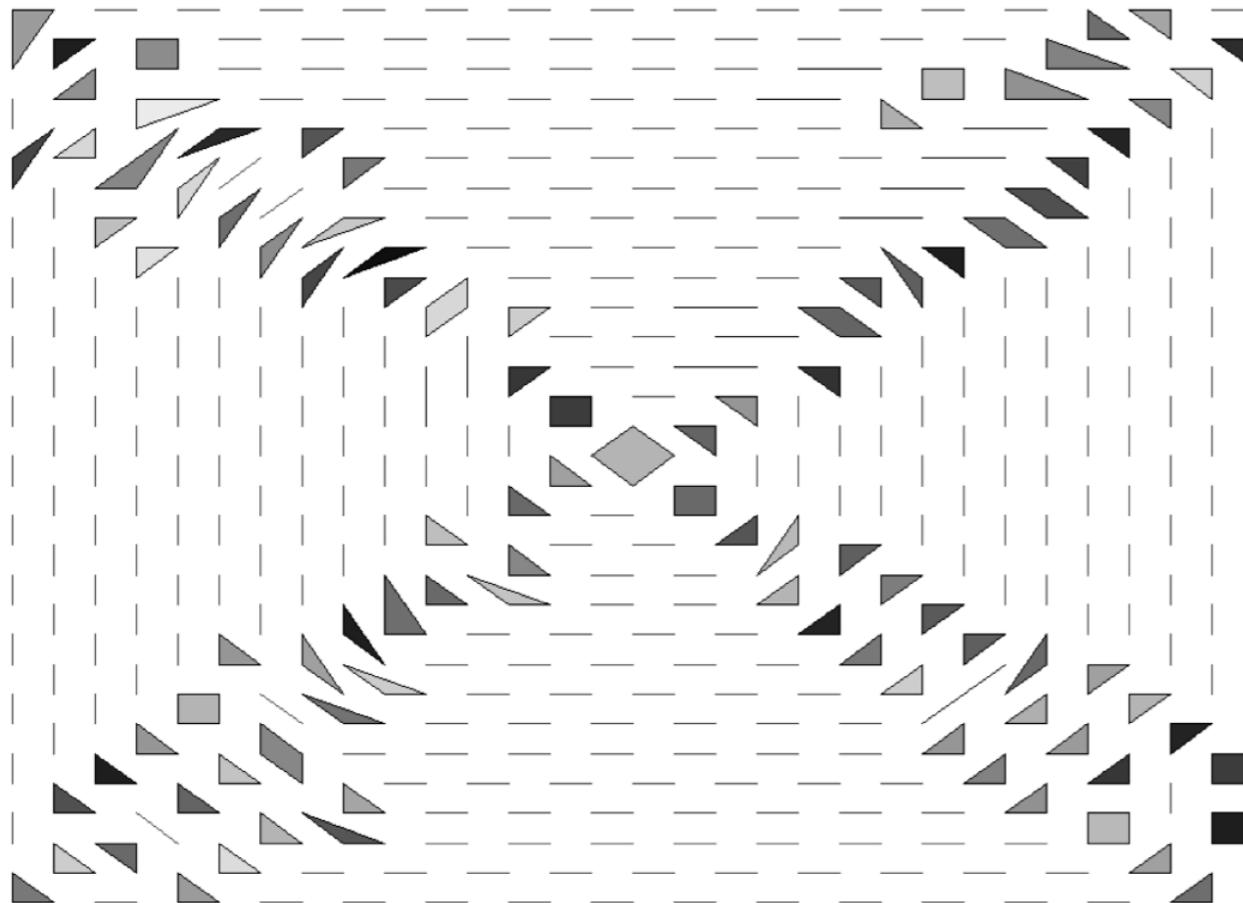


Evolution Measure

1. drop point source at a node

2. evolve point source with A

3. evaluate diffusivity at neighbors in comparison to known low energy



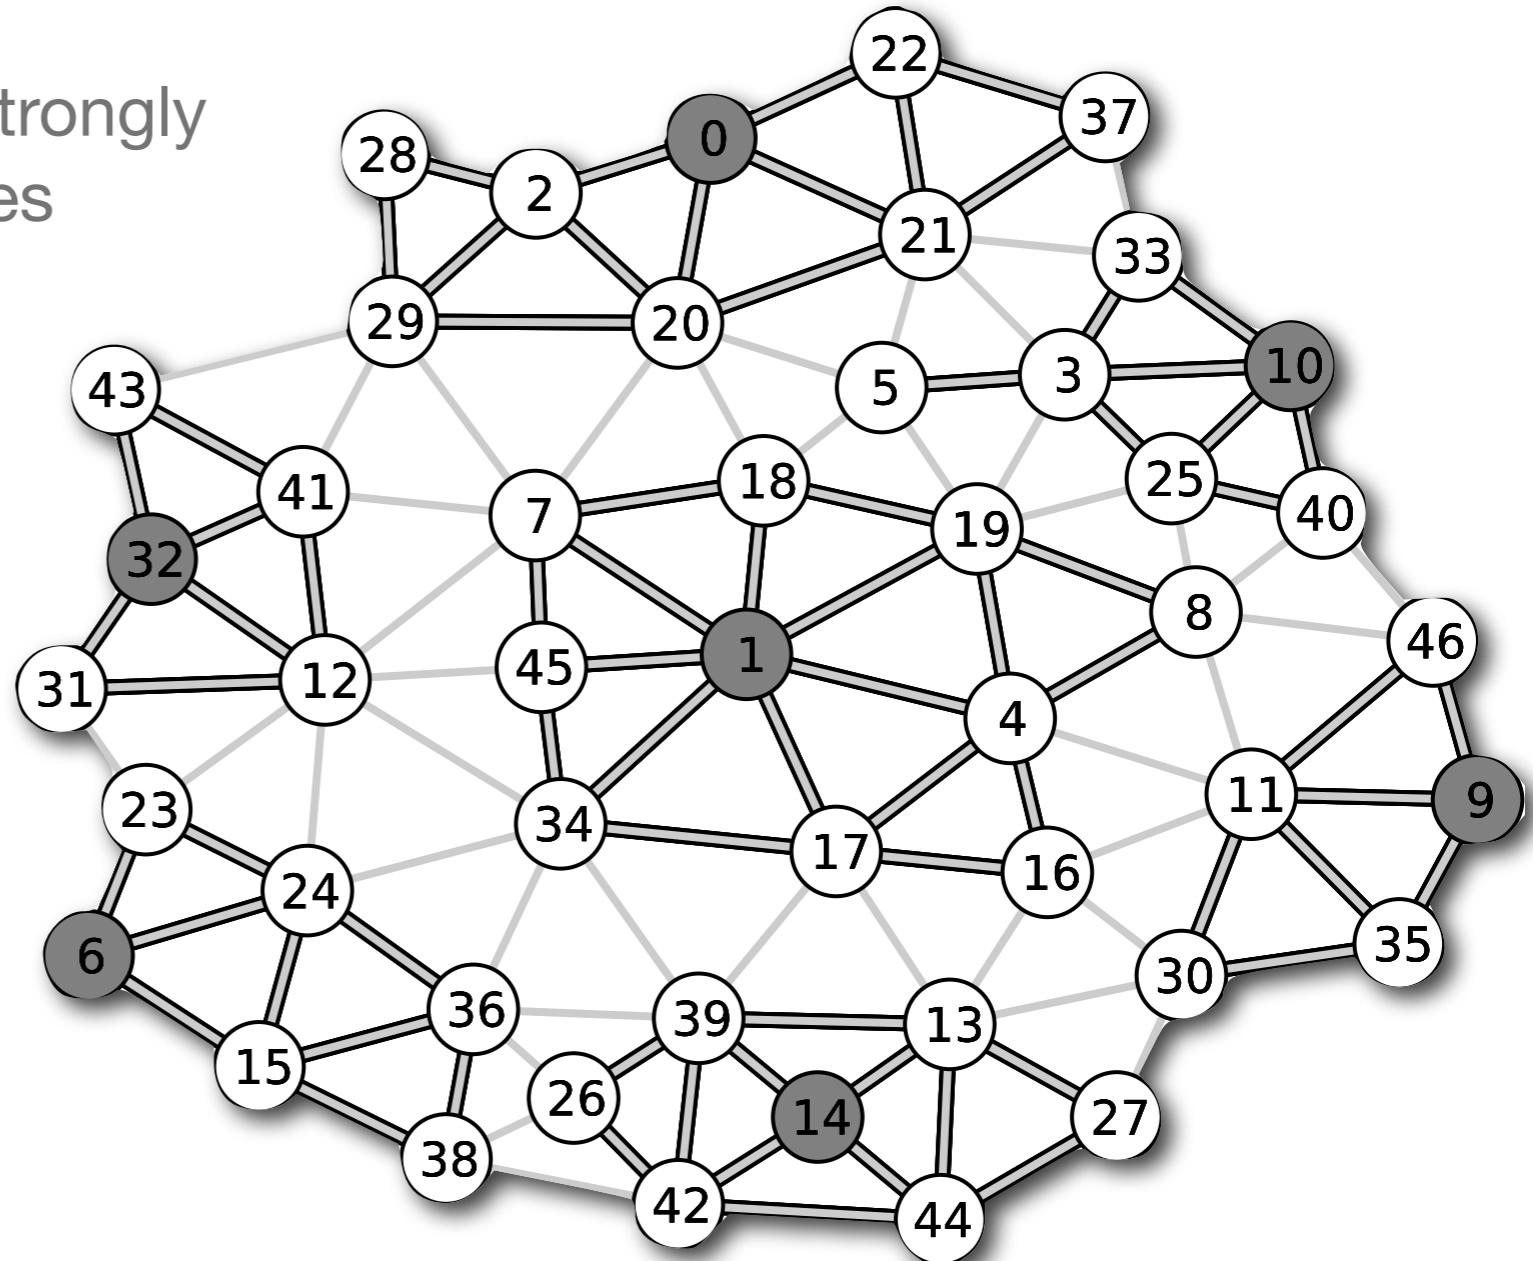
- efficient
- parameter insensitive
- Euler flow
- wave problems
- high-order
- discontinuous elements

Basic Graph Aggregation

- Metis
- Greedy: group collections of strongly connected unaggregated nodes

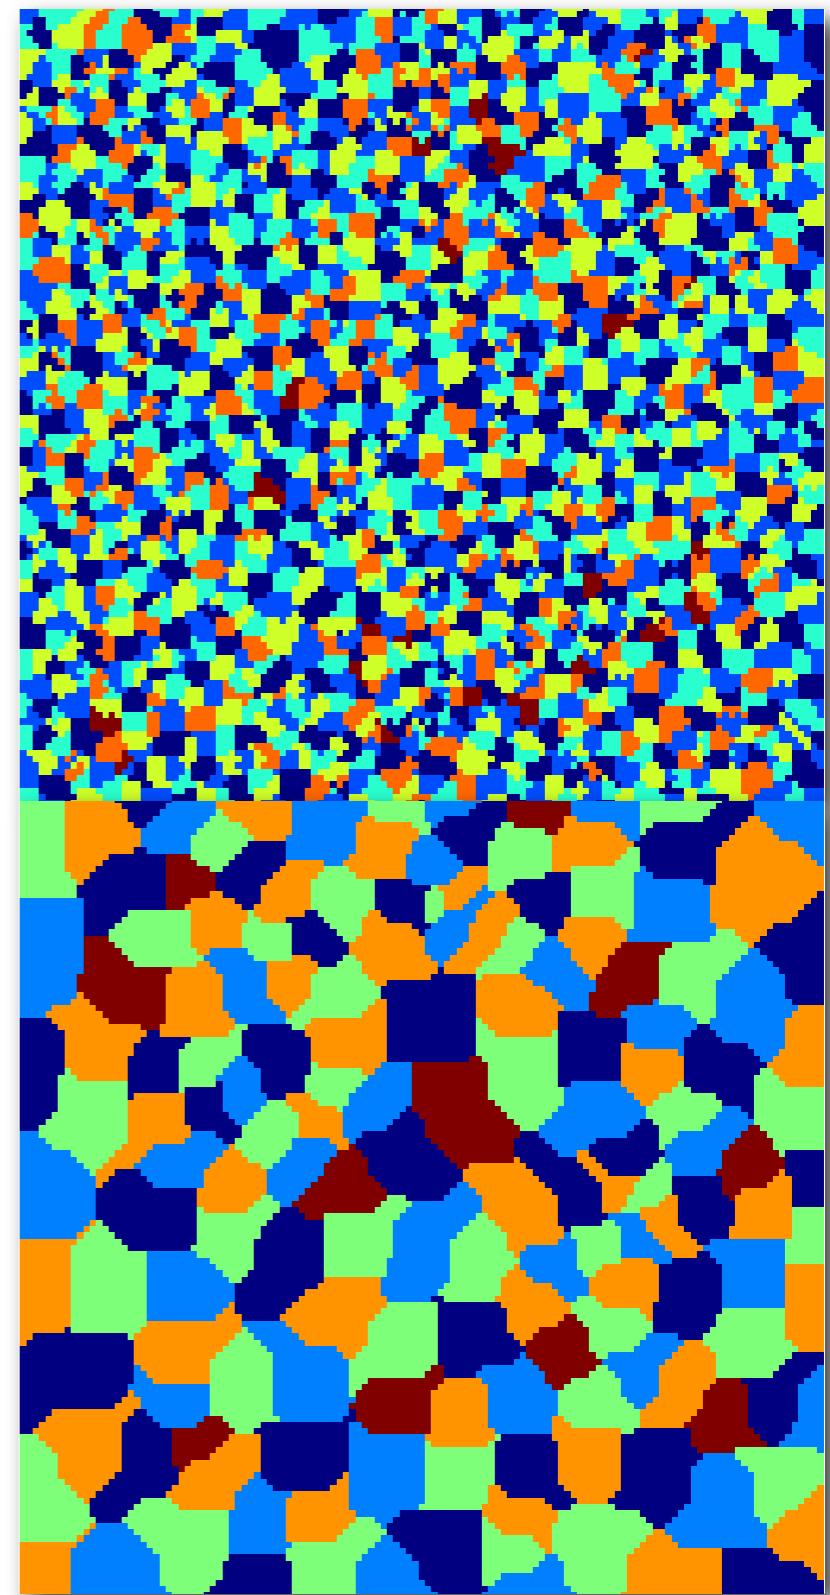
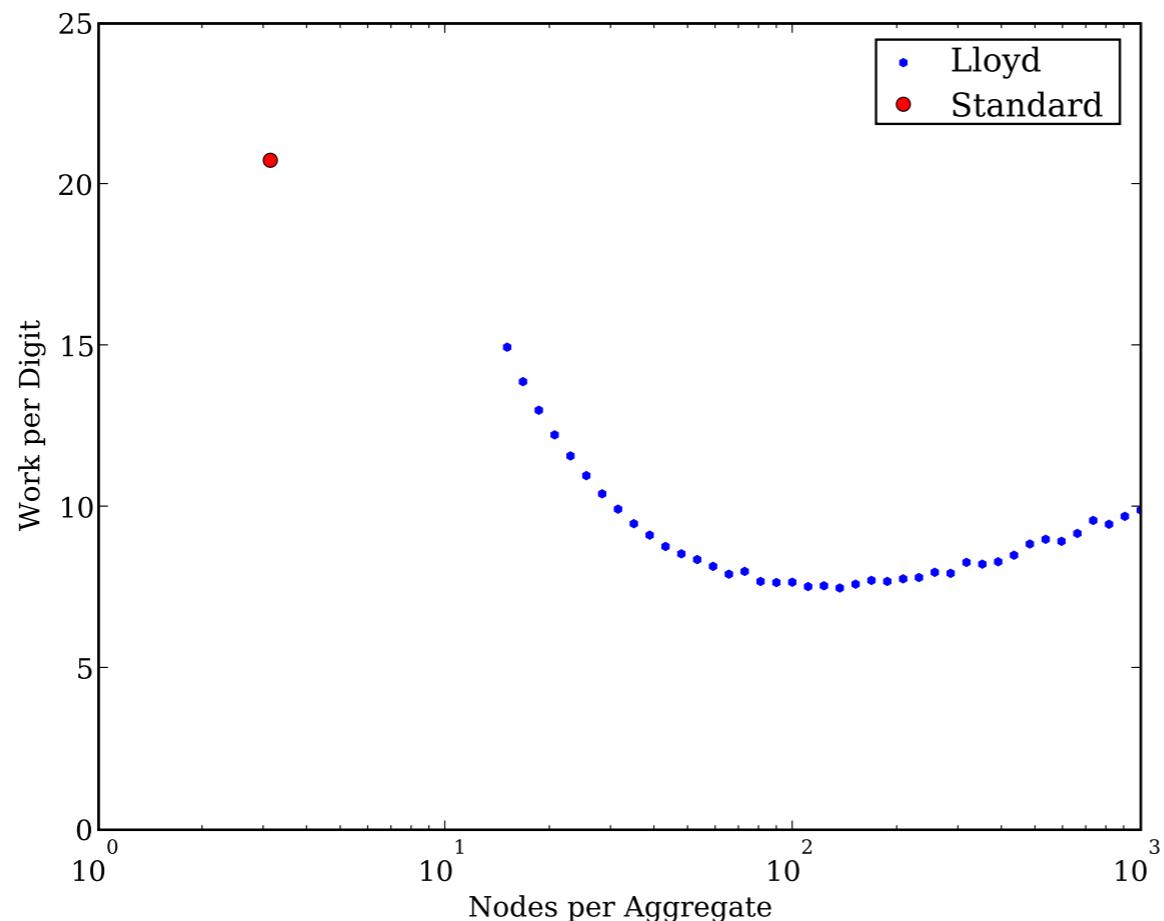
- Problems:

1. size fixed
2. sequential (greedy)

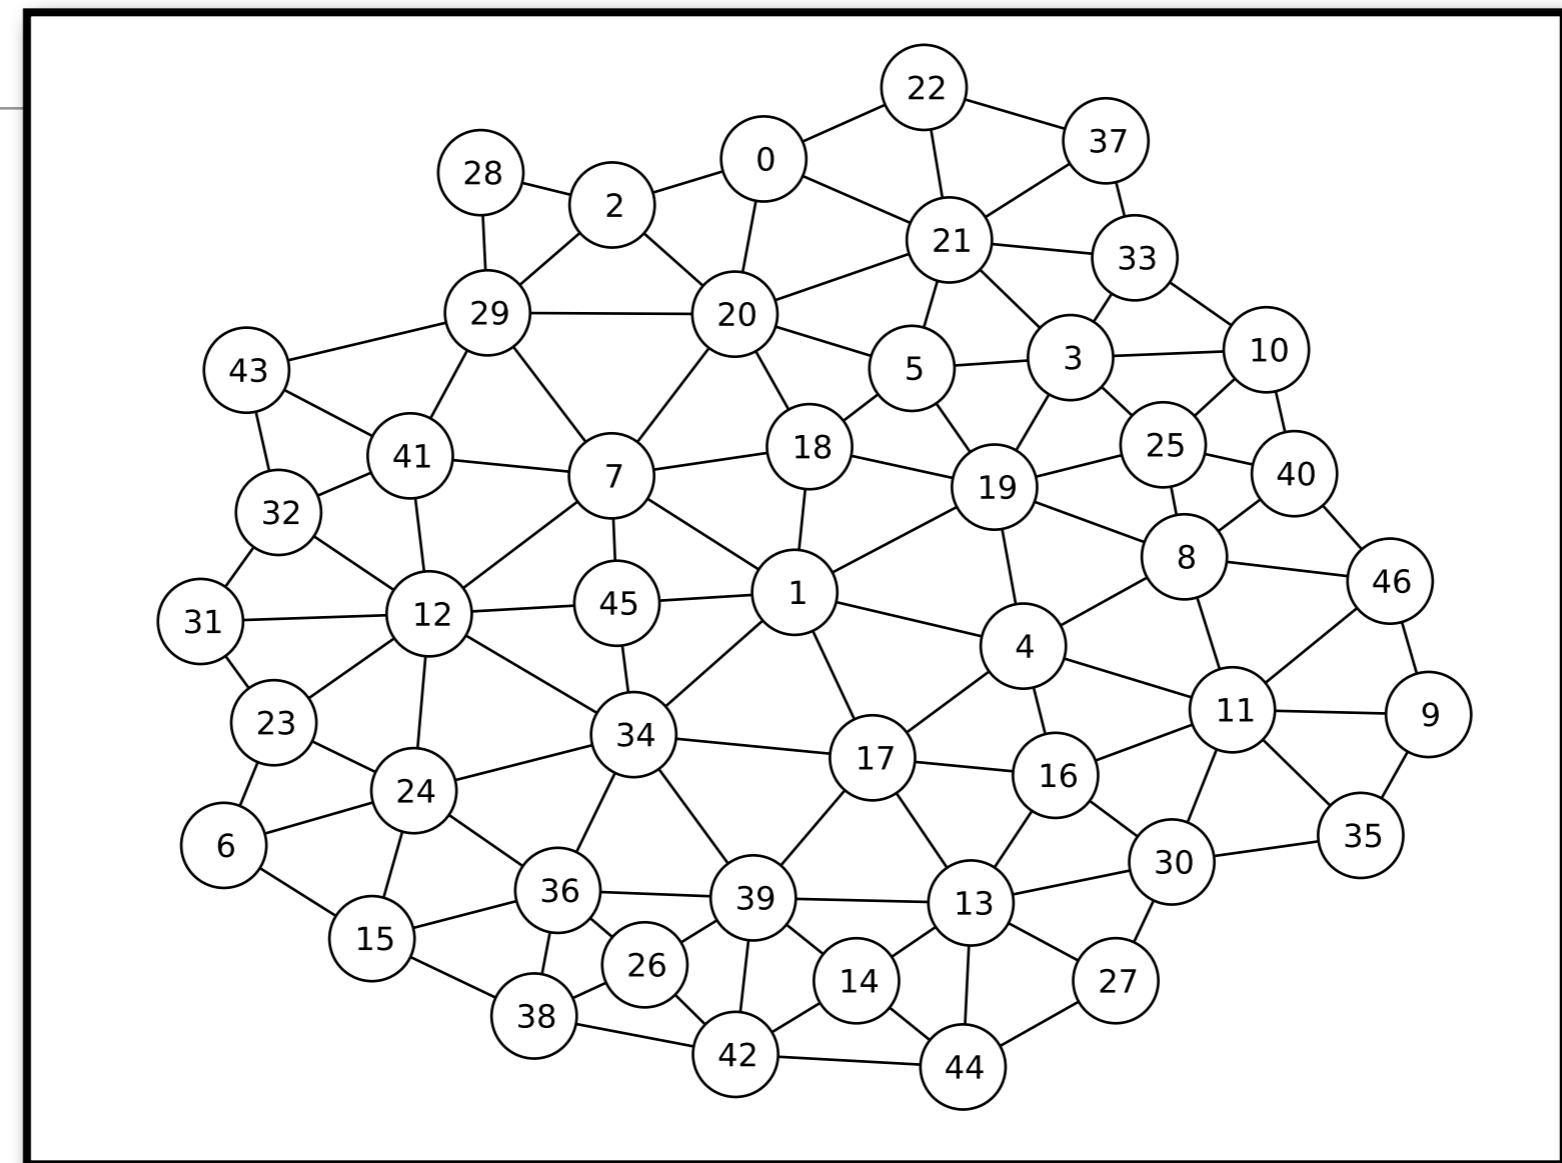


Flexible aggregation

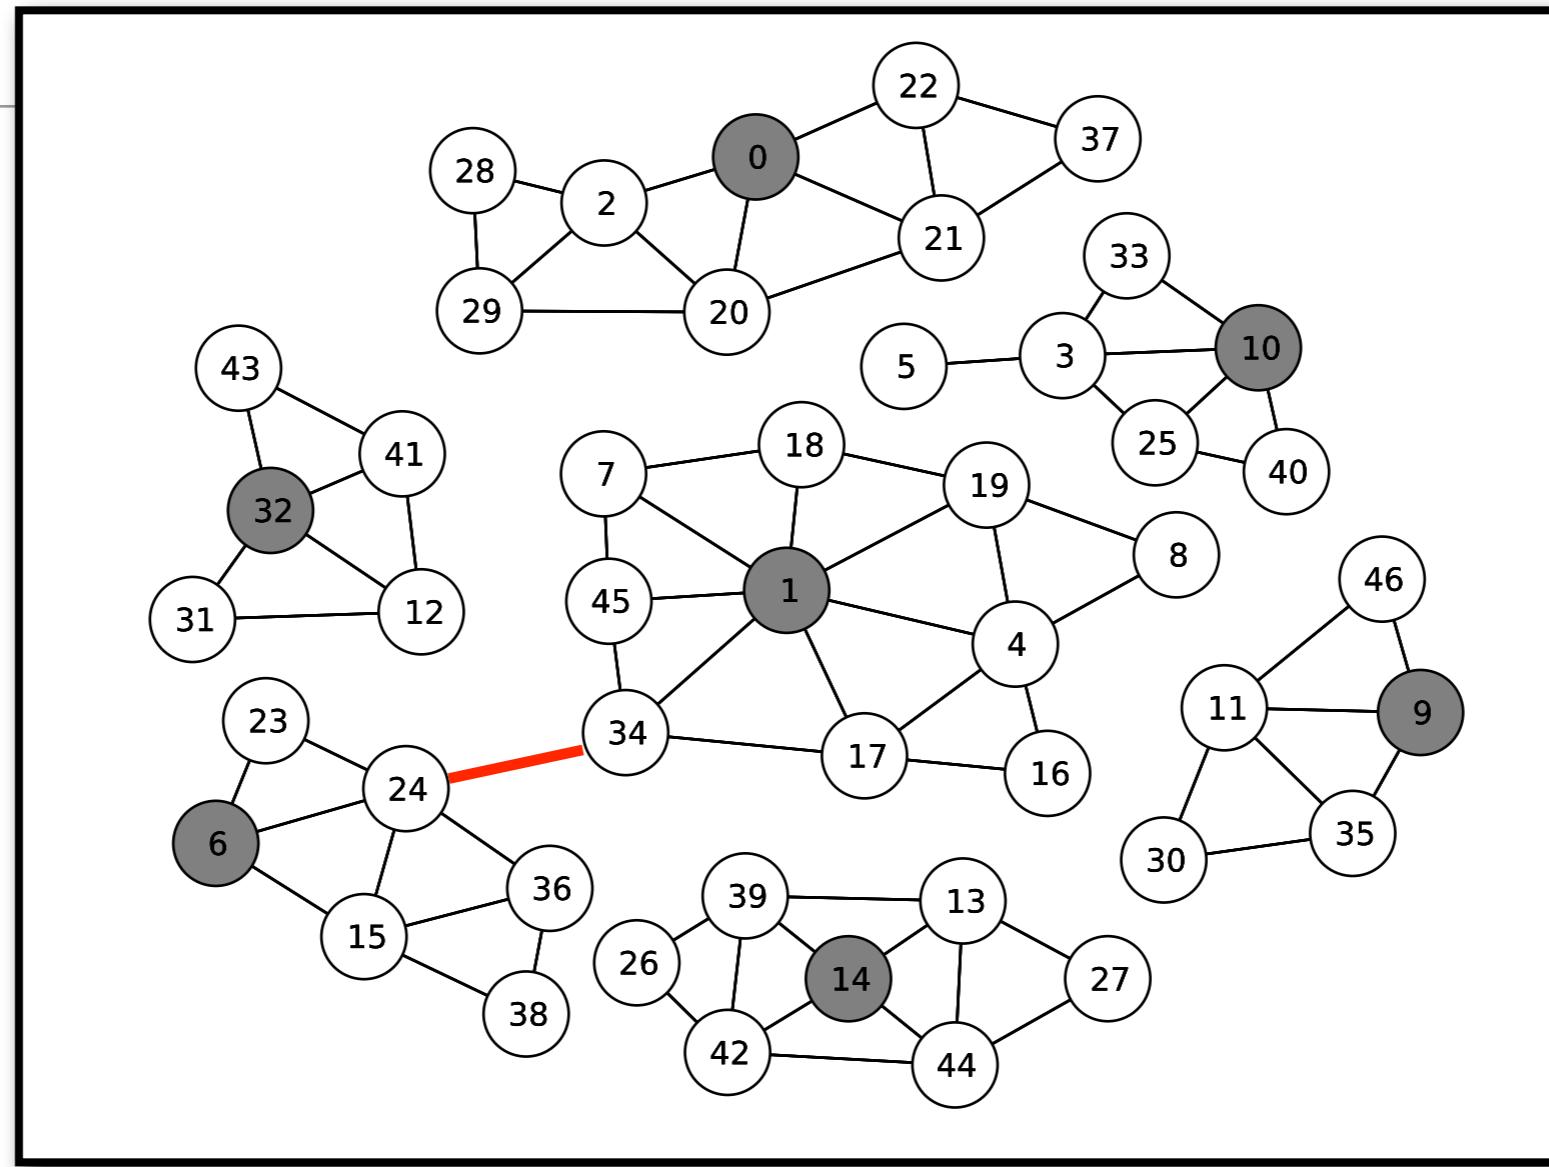
- two approaches:
 1. shortest-path: ability to tune bandwidth
 2. MIS(k): parallel coarse grids == serial coarse grids



MIS(2)



MIS(2)



- root nodes more than 2 edges apart ($>$ distance-2)
- an unaggregated node more than 2 edges from a root can become a root

independent
 } MIS(2)
 maximal

MIS(2)

1. Given a MIS(2) of {0,1} (N nodes)

2. prefix scan to enumerate

```
thrust::exclusive_scan(set, set + N, set, init)
```

3. communicate to neighbors
aggregate index (**SpMV**)

4. communicate index to unaggregated
neighbors (another **SpMV**)

} ~ std. first pass
}

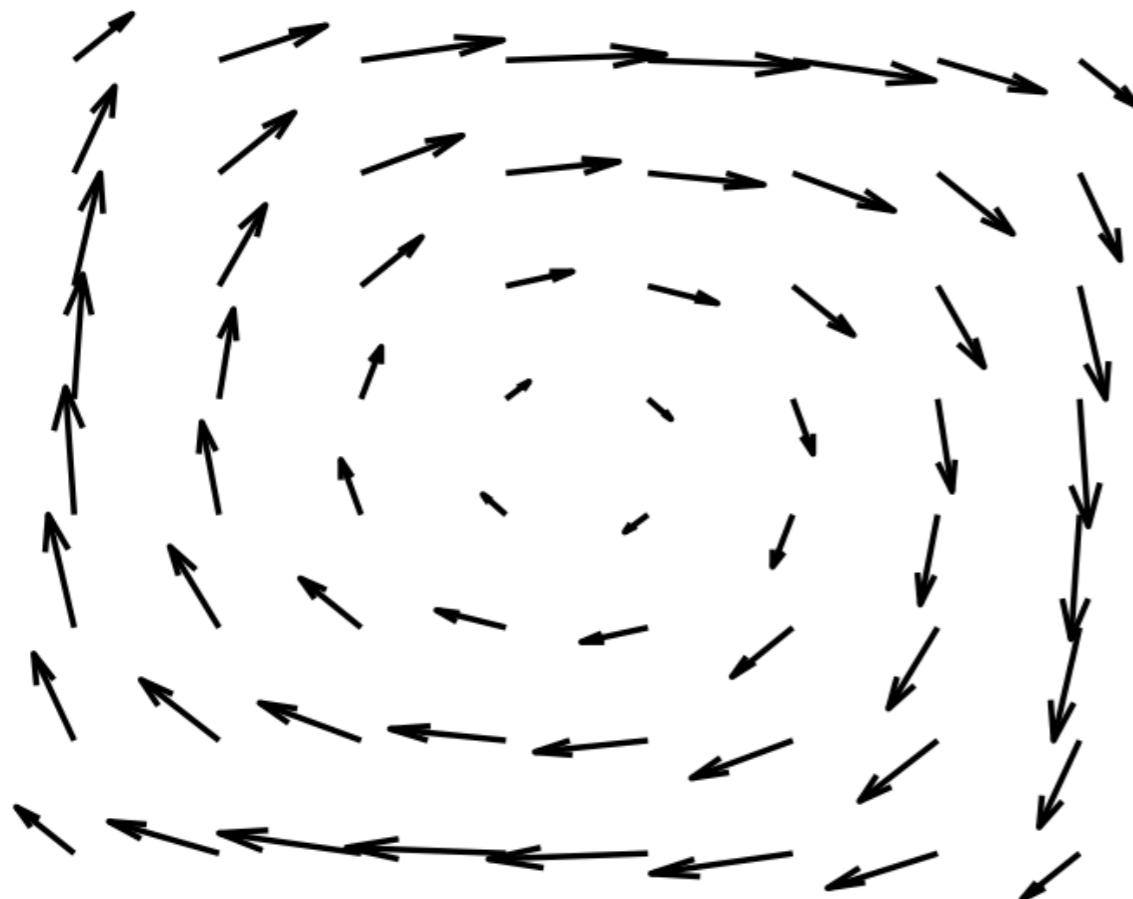
} ~ std. second pass

- Generalizes to MIS(k)
- allows for variable coarsening

*** Olson, Schroder, Tuminaro, *A general interpolation strategy for algebraic multigrid using energy-minimization*, 2010.

optimizing energy

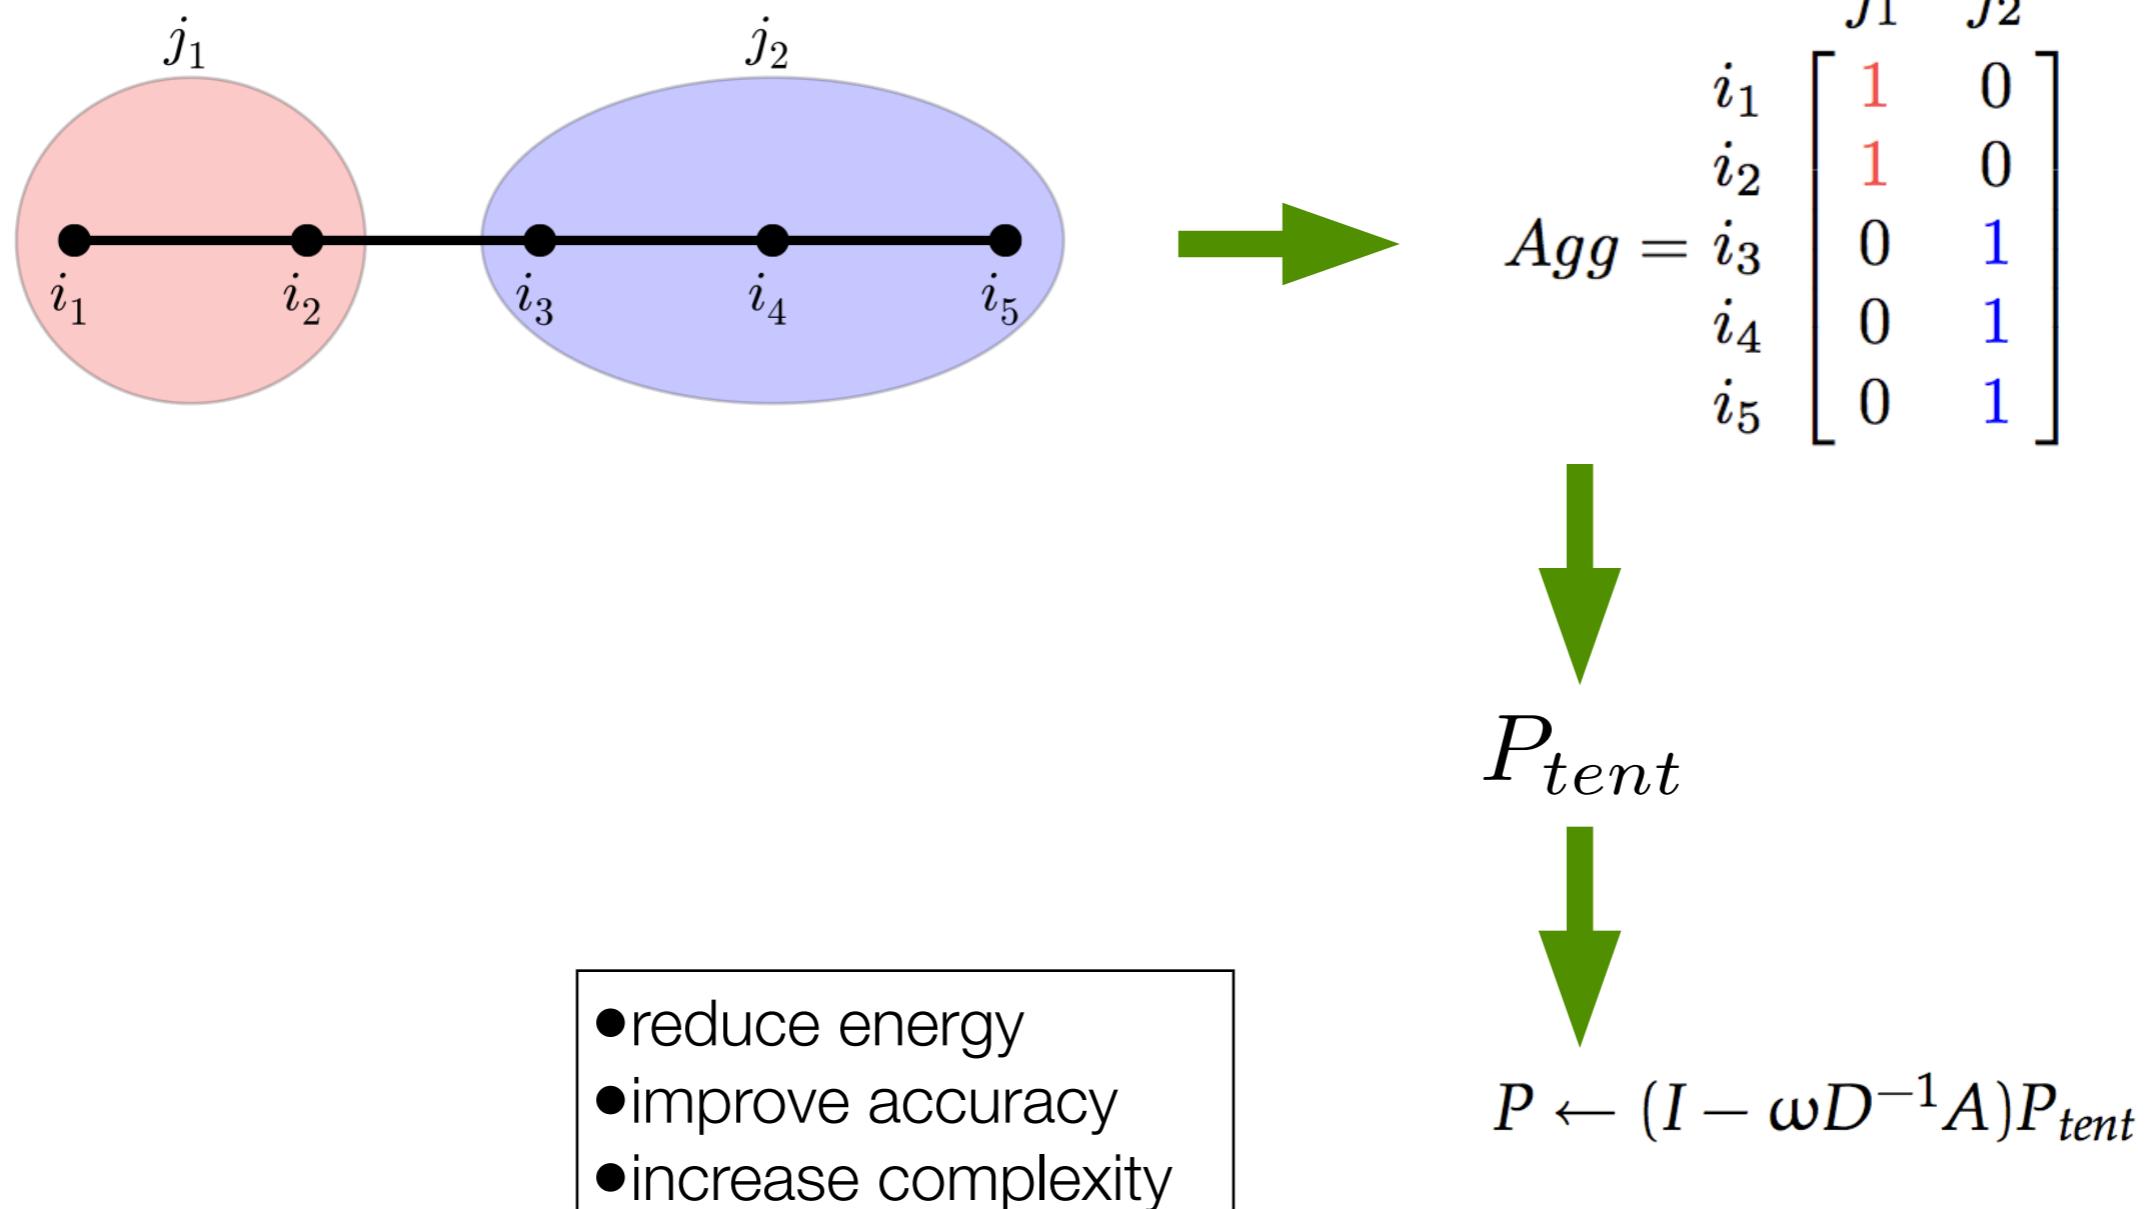
- P should have low energy
(low A -norm or A^*A -norm)
 1. determine sparsity pattern
 2. minimize energy column-wise (parallel)



	std.	opt.
1/64	>150	24
1/128	>150	28
$\frac{h}{1/256}$	>150	33
1/512	>150	33

Interpolation: standard approach

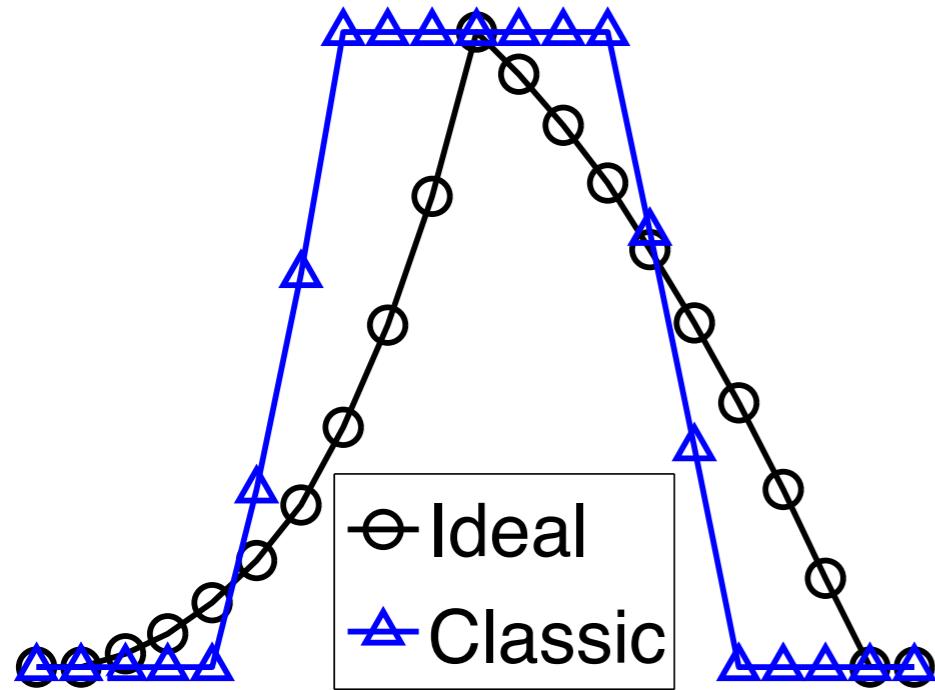
- Set the sparsity pattern from aggregation



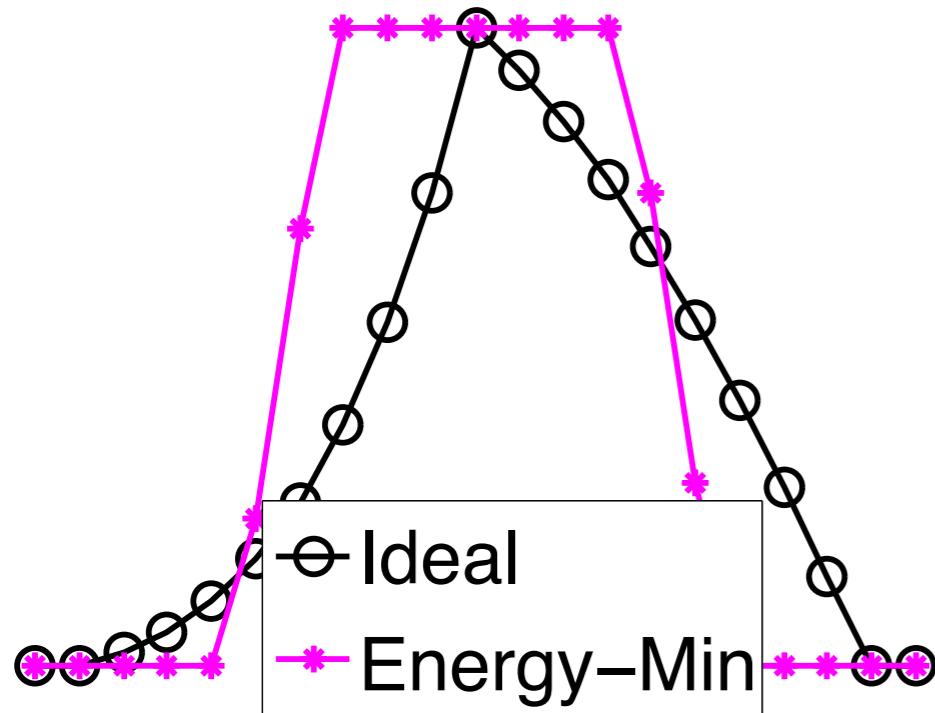
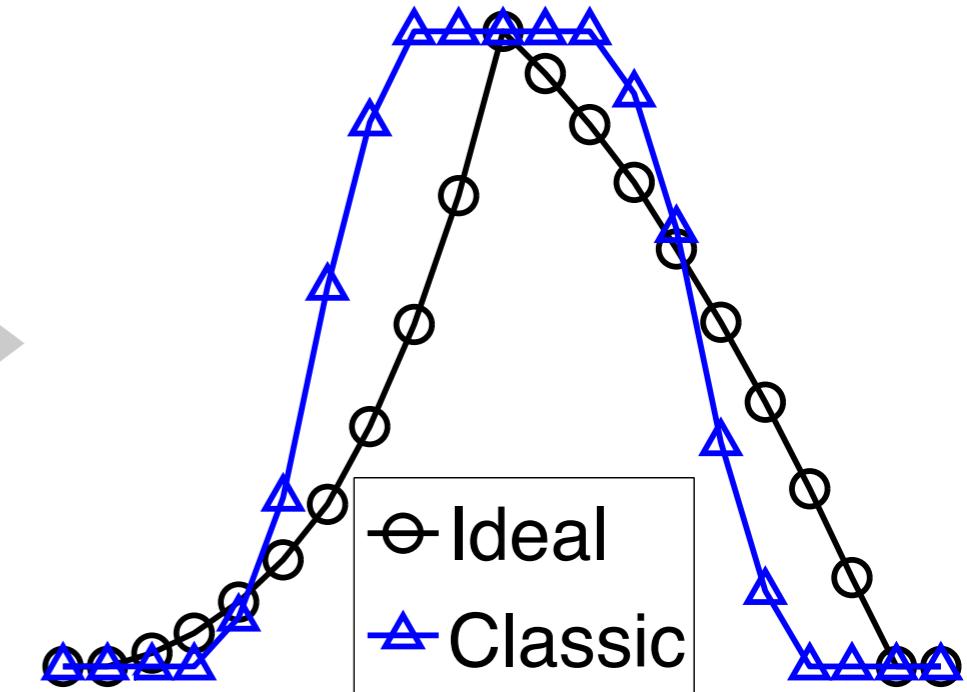
Idea #4

Range of Interpolation

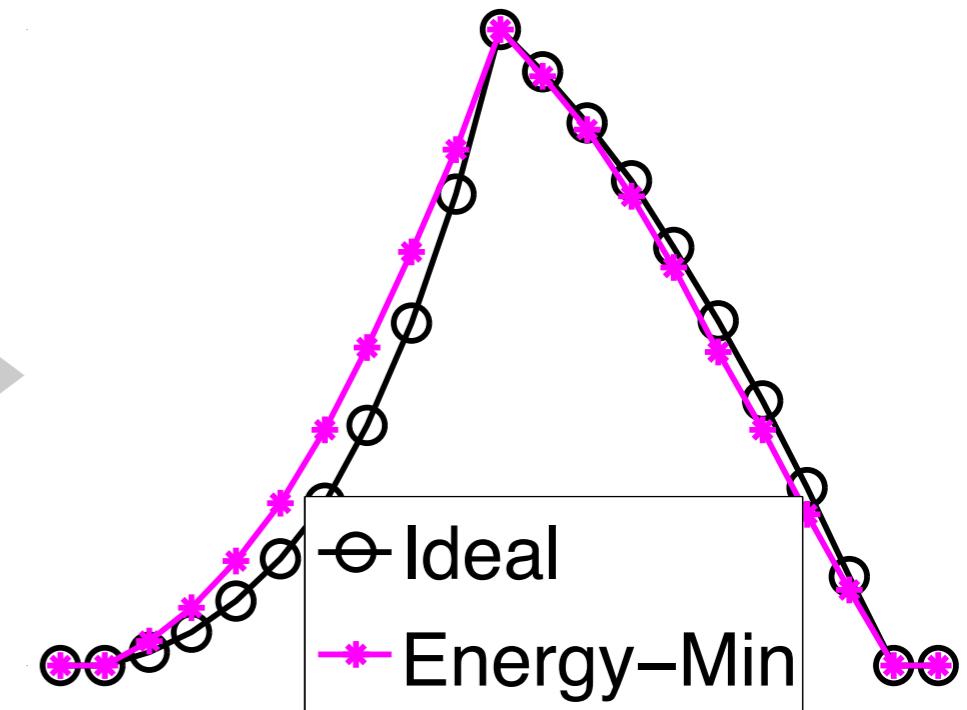
$$-u_{xx} + \sin(x)u_x = f$$



old



new



Toward General Interpolation

- Want P so that $u_{low} \in \mathcal{R}(P)$

1. Grow and fix sparsity pattern as $S^k P_{tent}$

2. Minimize residual of

$$AP_j = 0 \quad \text{for each column } j$$

3. Constraint the minimization with

$$PB_c = B$$

Toward General Interpolation

- Hermitian (and positive definite): use CG

$$AP_j = 0 \Leftrightarrow \min \|P_j\|_A$$

$$R = P^*$$

- Non-Hermitian: use GMRES

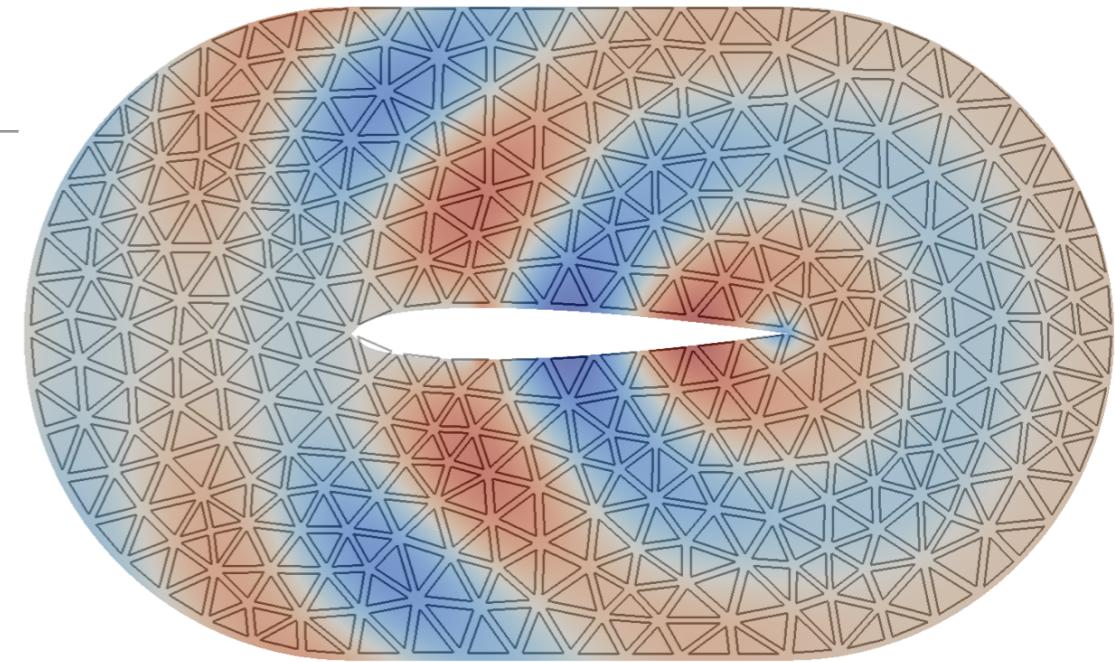
$$AP_j = 0 \Leftrightarrow \min \|P_j\|_{A^* A}$$

$$A^* R_j^* = 0 \Leftrightarrow \min \|R_j^*\|_{AA^*}$$

- Range of interpolation targets “right” low-energy
- Range of restriction* targets “left” low-energy
- Cost is comparable to that of standard smoothing

Wave problems (Helmholtz)

- AMG problem: standard “low energy” $e^{i\omega x}$ modes (constant) break Nyquist rate
- answer: introduce wave modes at all levels

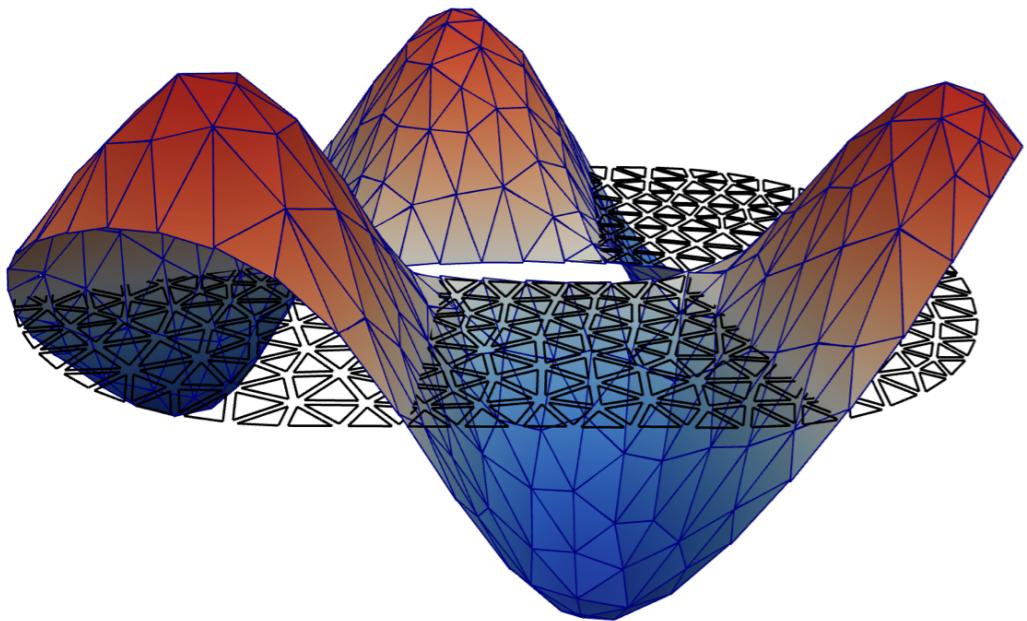


multiple wave modes: $\rightarrow e^{i\omega \cos(\theta)x + \sin(\theta)y}$

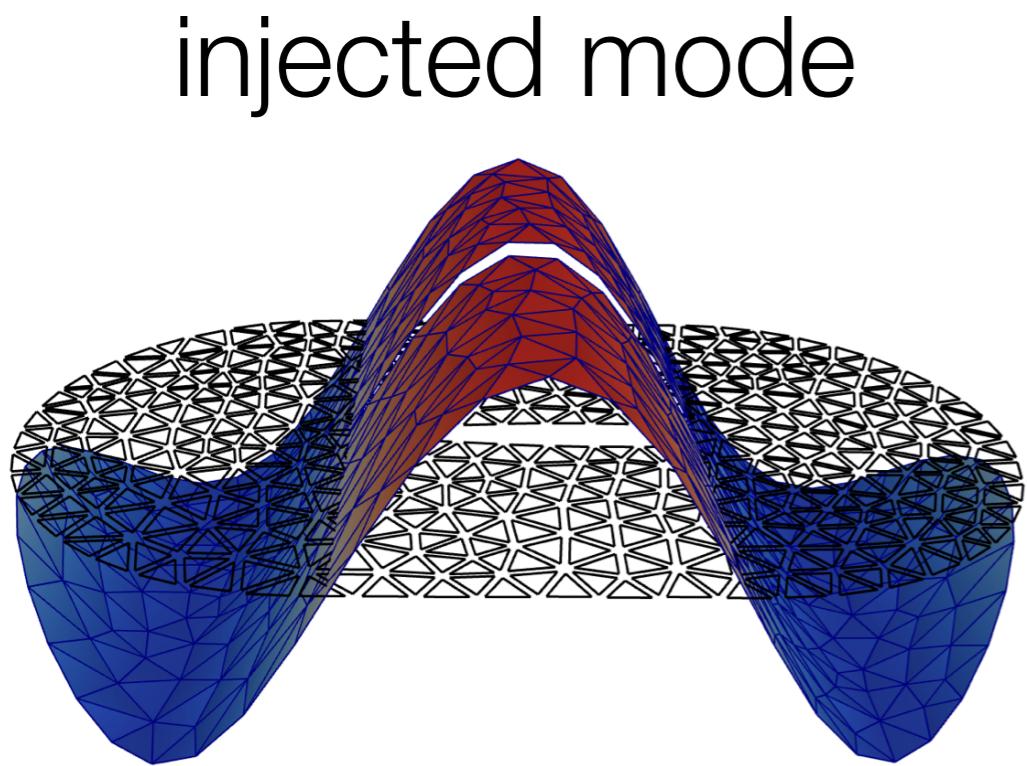
infuse wave modes: \rightarrow
 level 1: $\theta = 0, \frac{\pi}{2}$
 level 2: $\theta = 0, \frac{\pi}{8}, \frac{3\pi}{8}, \frac{5\pi}{8}, \frac{7\pi}{8}$

adapt to grid through relaxation: $\rightarrow Ae^{i\omega \mathbf{x}} = 0$

optimize interpolation: $\rightarrow \|P\|_{A^* A}$



left singular vector



injected mode

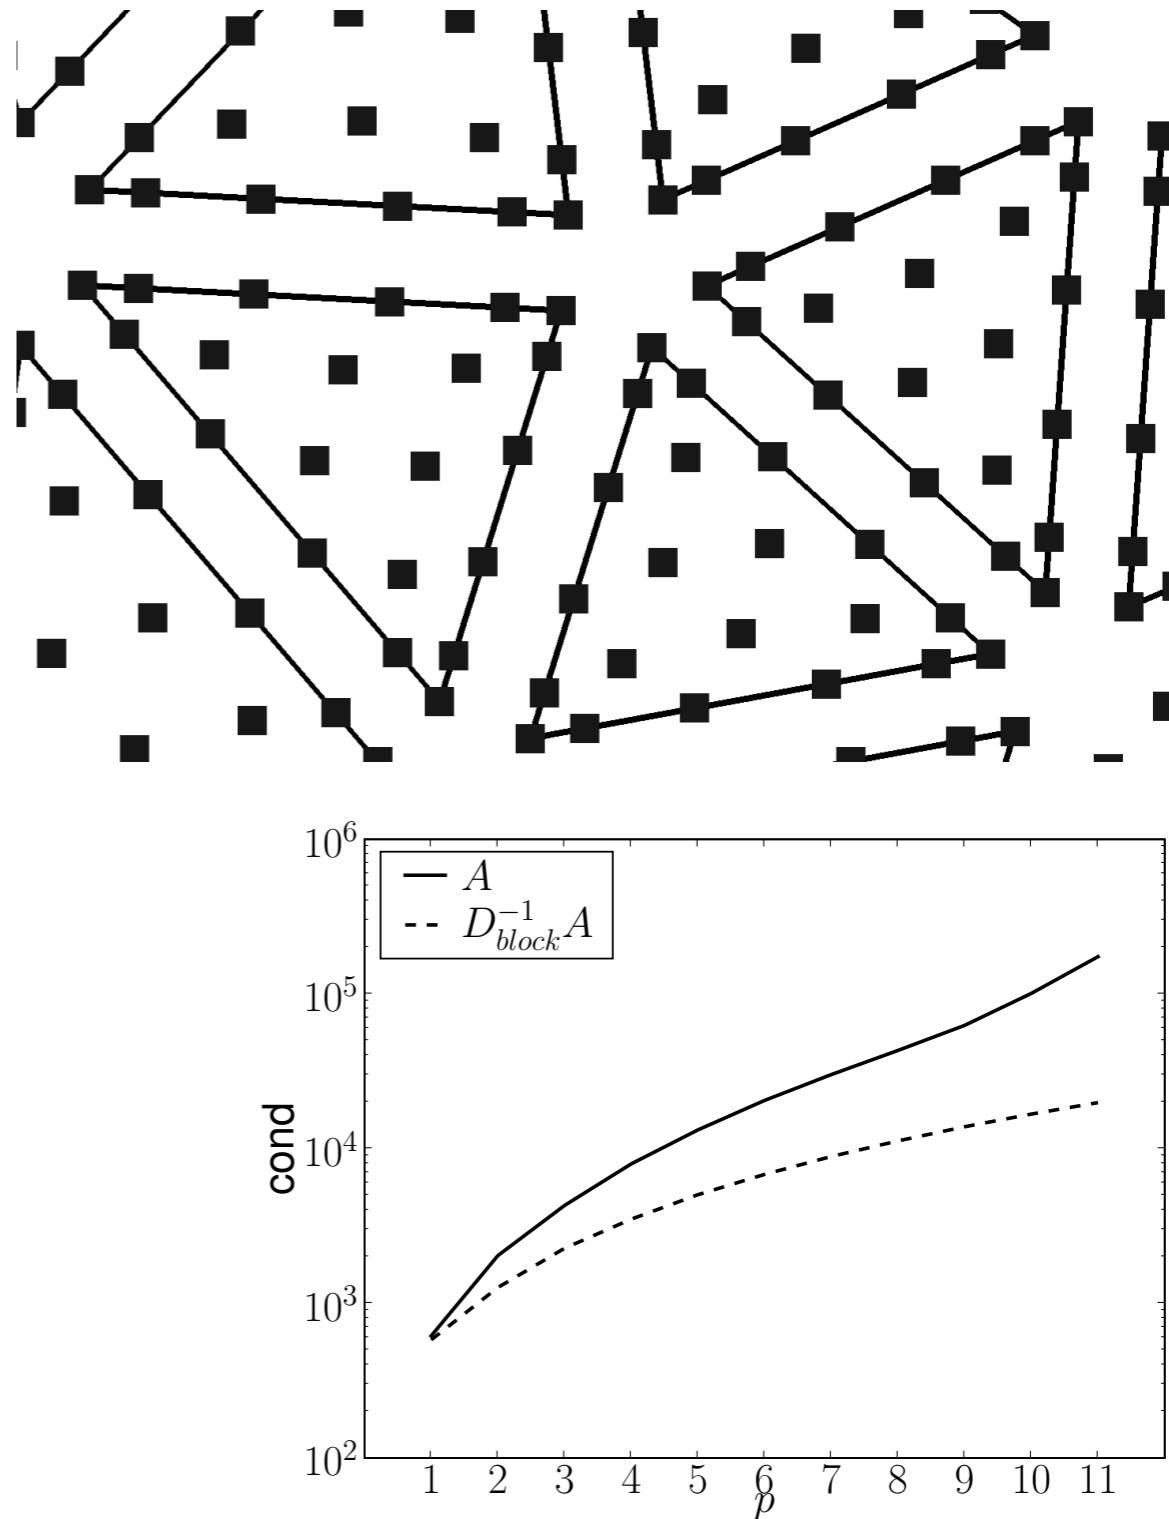
*** Olson, Schroder, *A smoothed aggregation multigrid method for Helmholtz problems*, 2010.

ppw	h	$h/2$	$h/4$	$h/8$	$h/12$
5.0	9	11	15	25	54
10.0	9	11	11	11	14
30.0	9	10	10	11	11
90.0	9	10	11	11	11
All	4	15	17	21	22

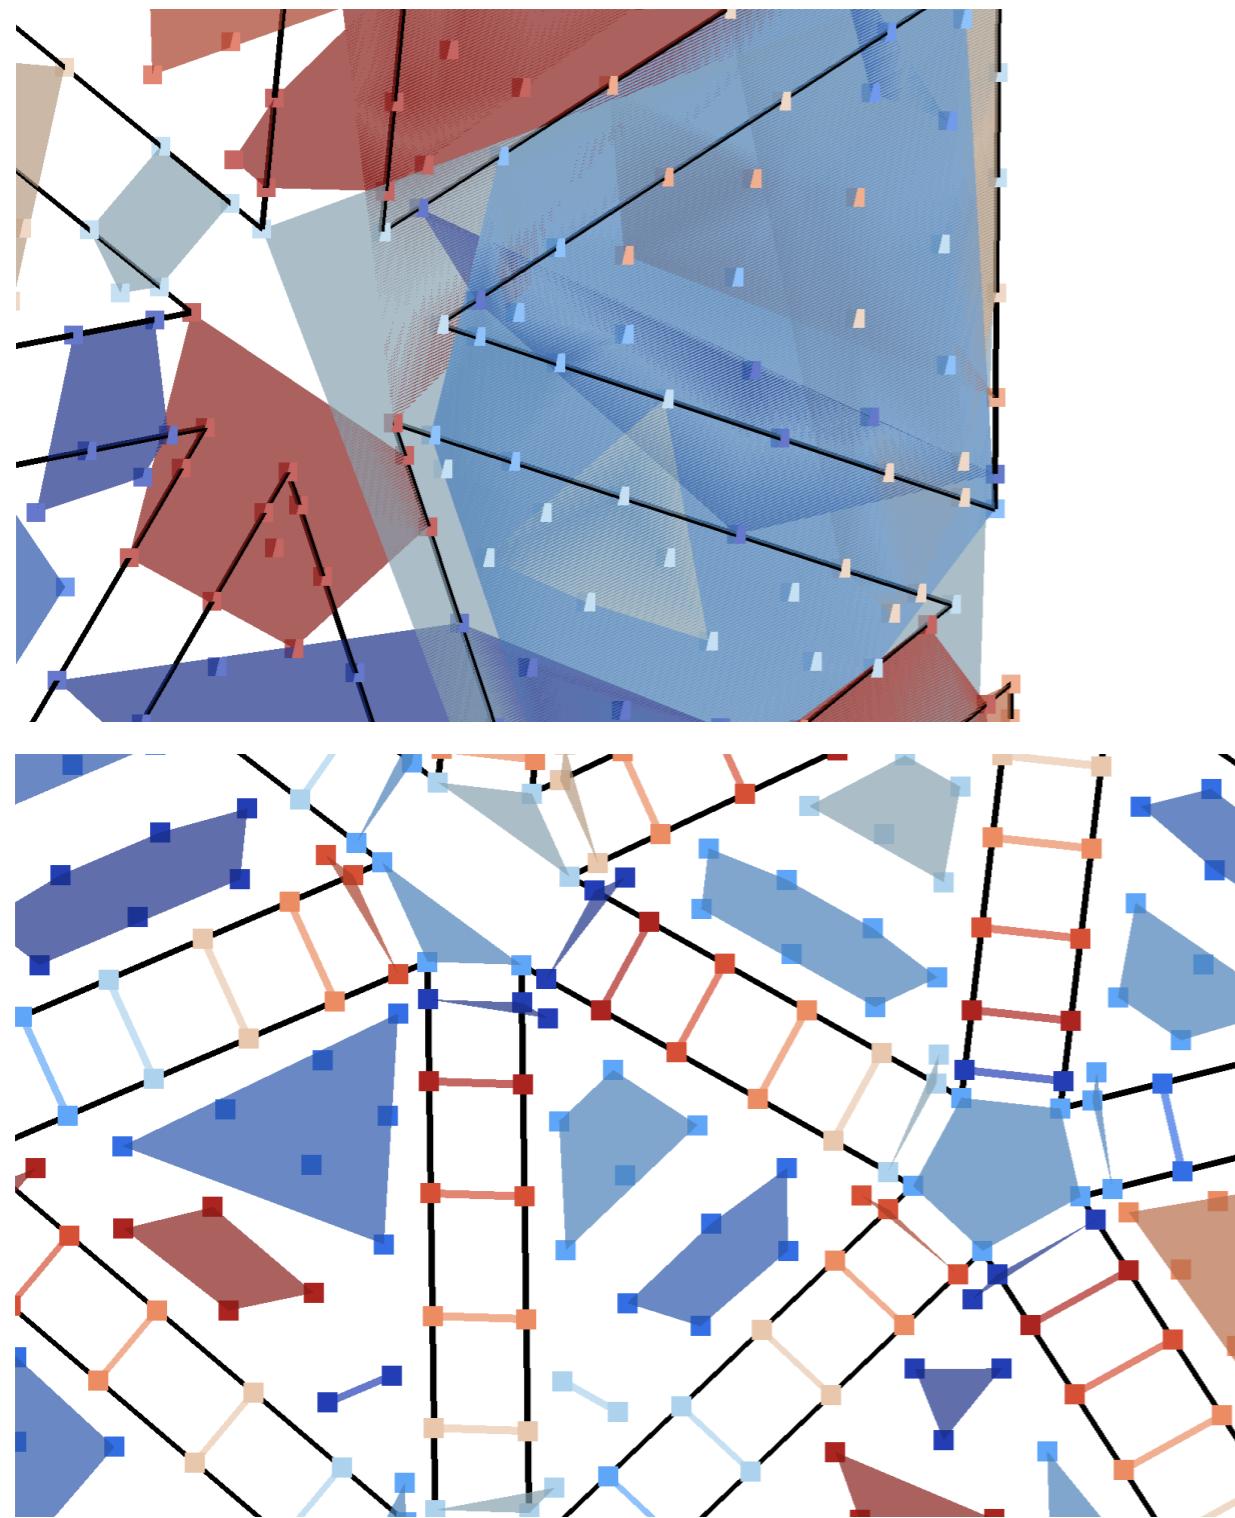
pgmres iterations, complexity

High-order discontinuous Galerkin

- different types of d.o.f.
- loss of locality
- increase in condition number
- most heuristics in AMG break down

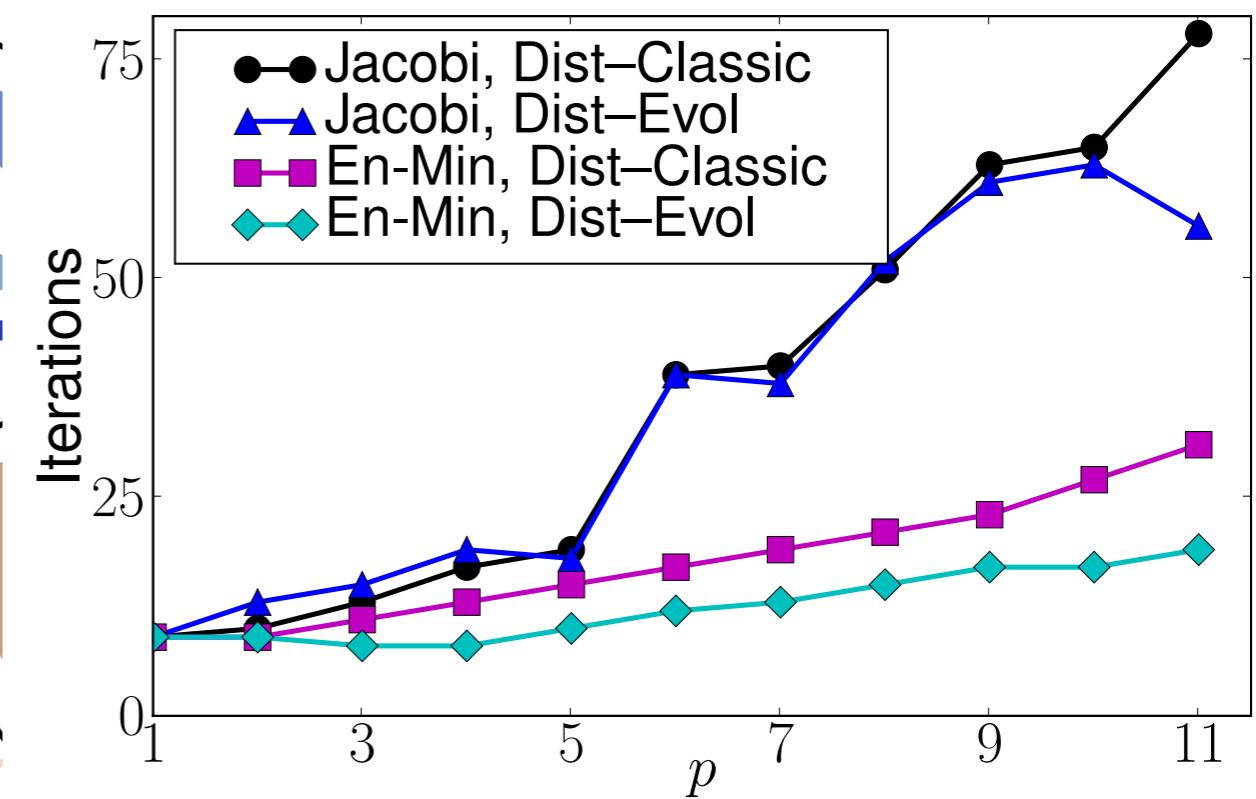


High-order discontinuous Galerkin: Poisson



key ingredients:

- conforming aggregations step
- adapt the near null space
- optimal interpolation



Acceleration Motivation 1: potential

1M d.o.f.
10 nonzero neighbors
1.5 cycle complexity
15M computations

2 for a $V(1,1)$ -cycle
20 iterations
600M for solution

16 bytes for (data, row, col) (8,4,4)
10 Gbytes of data to process

125 Gbyte/sec SpMV on the GPU?
~ 0.08 sec

25 Gbyte/sec SpMV on host?
~ 0.4 sec

Acceleration Motivation 2: software + hardware

real acceleration units: dedicated computation, double precision, high speeds

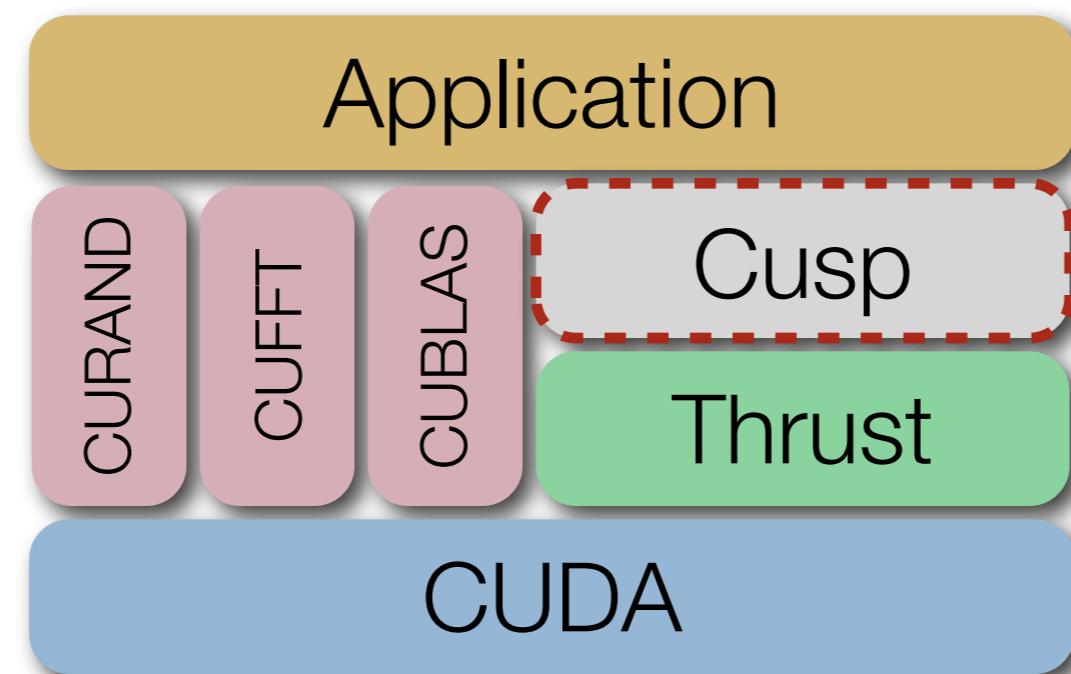
useable software: CUDA + Thrust + Cusp

AMG “asks” for acceleration:

- ✓ adaptive
- ✓ multiple candidates, cycles
- ✓ less/more aggressive coarsening
- ✓ multiple RHS

- fast development
- low overhead
- open source

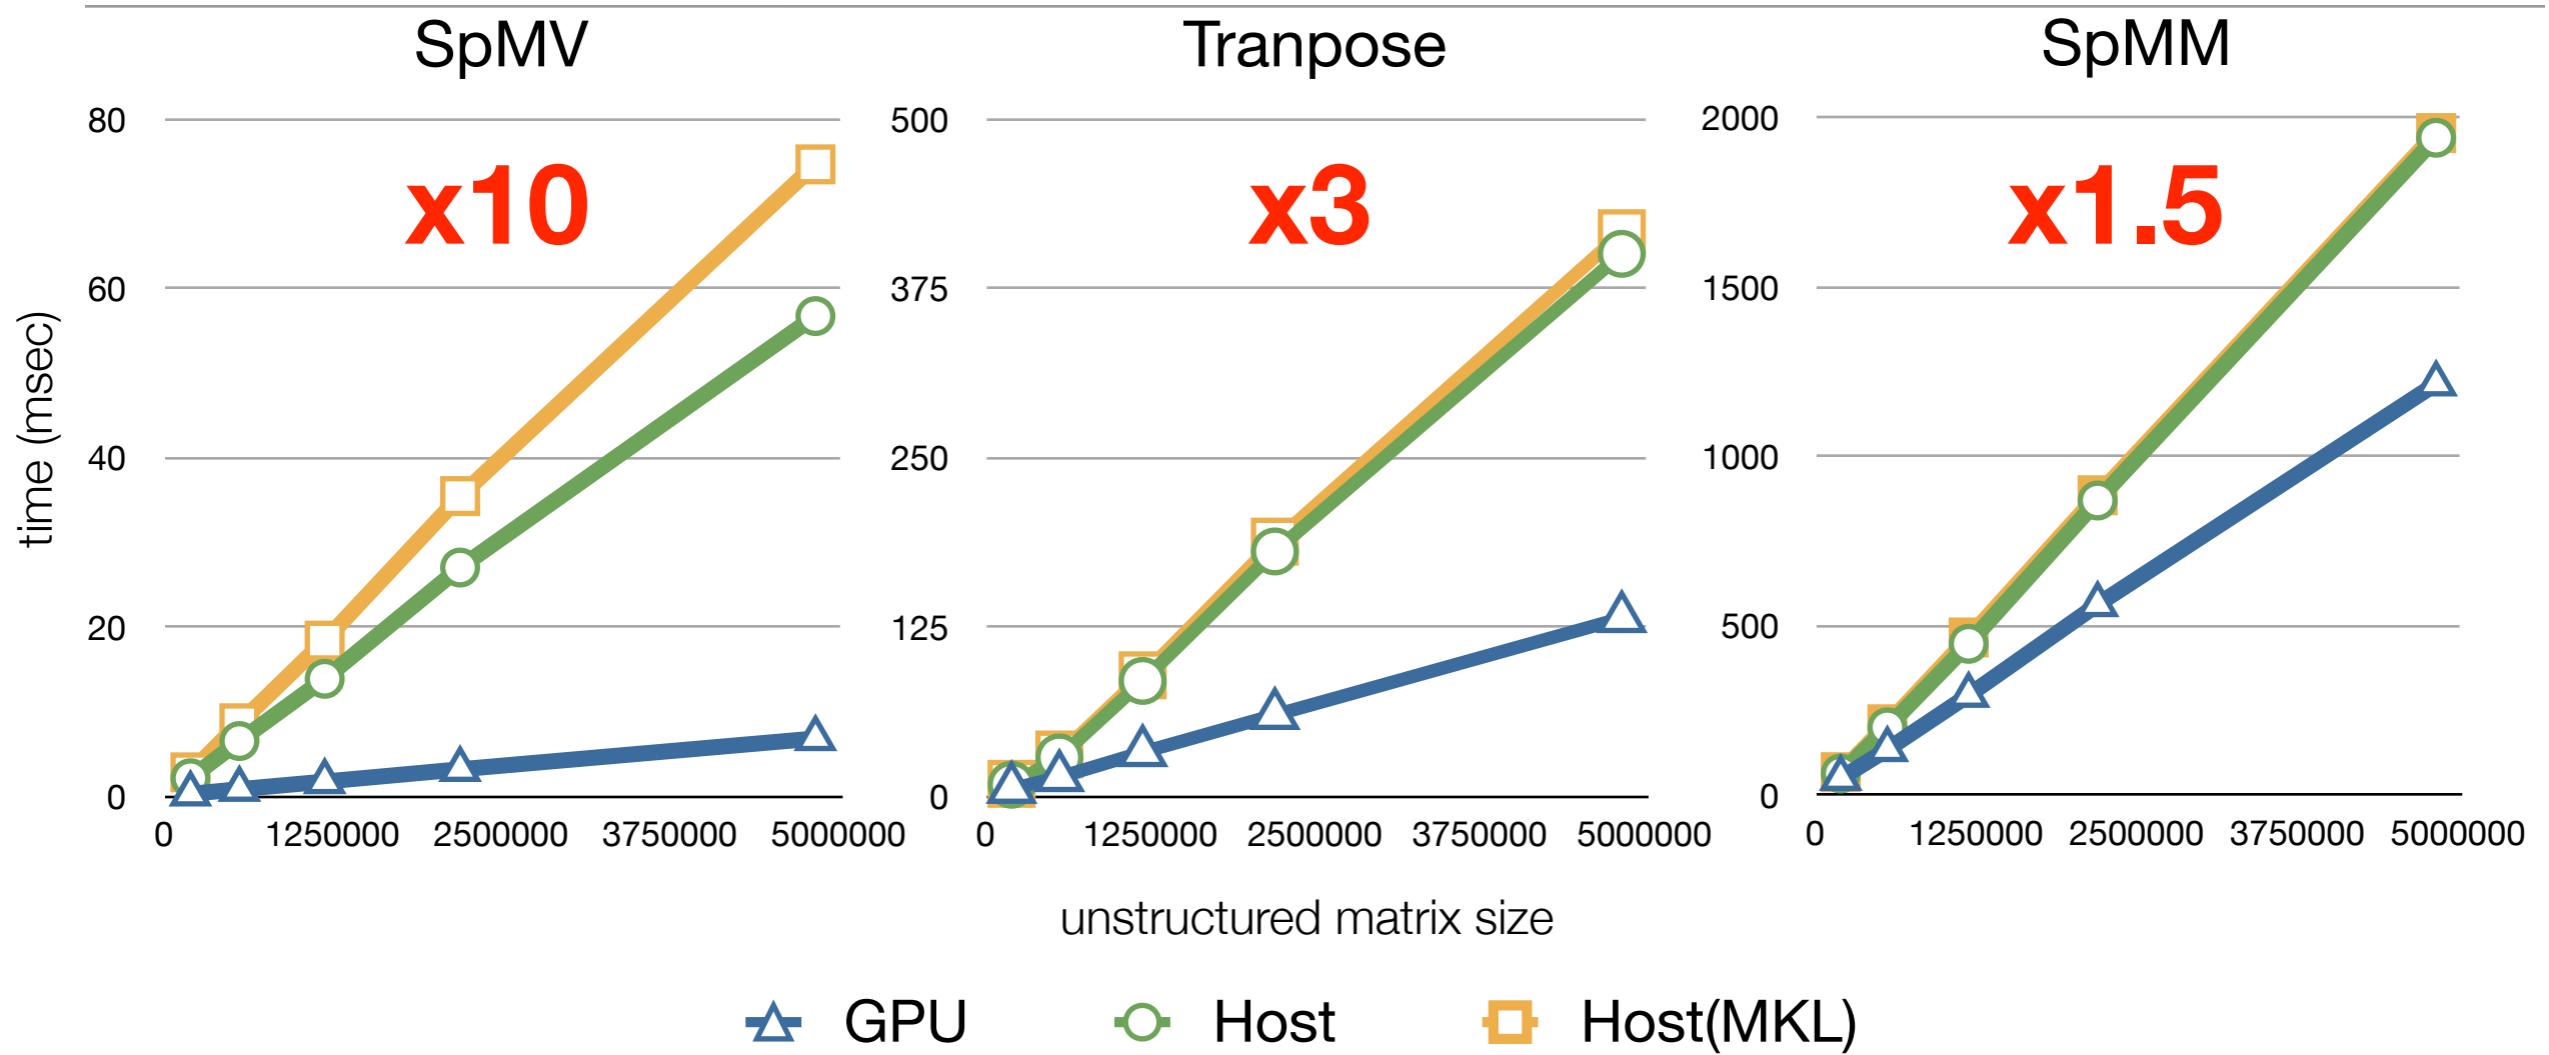
*** Bell, Dalton, Olson, *Exposing fine-grained parallelism in algebraic multigrid*, 2011.



Method:

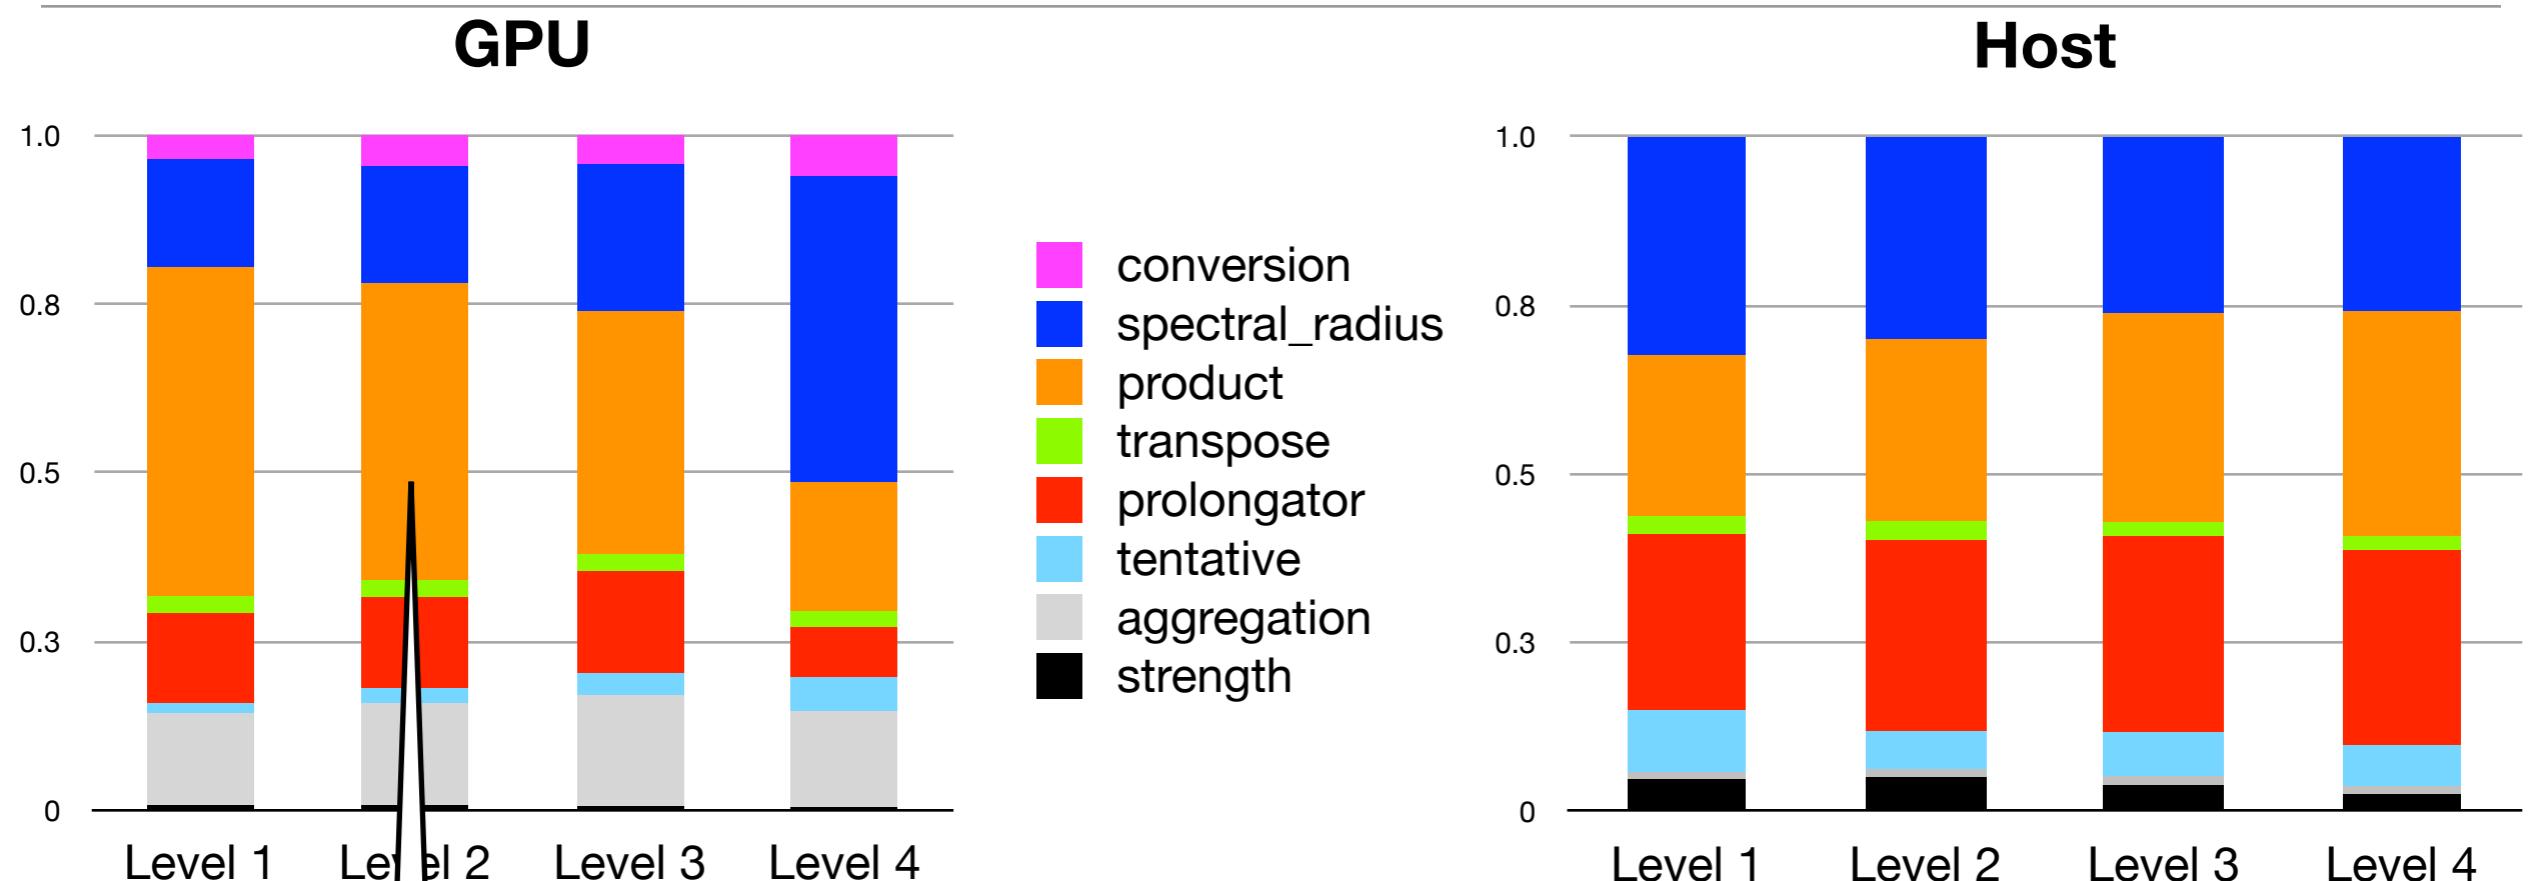
- expose fine-grained parallelism
- utilize fast kernels (gather, scatter, scans, sort, etc)

Setup kernels: SpMM, Transpose, SpMV



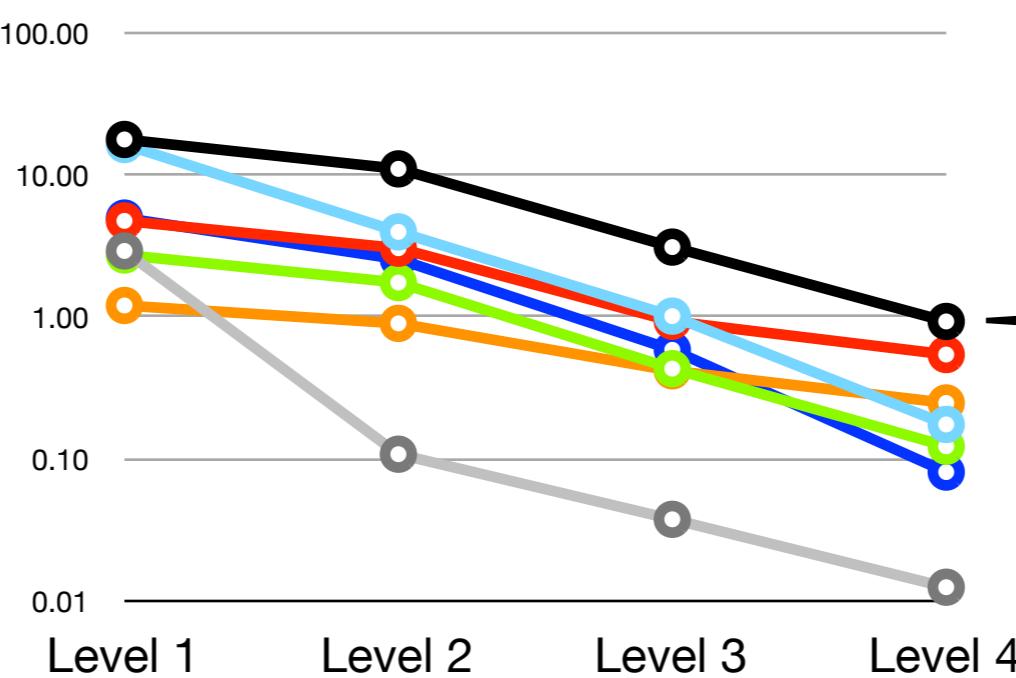
size	A*B	BT*(A*B)	BT*A*B	Transpose	SpMV
speed up	5E+06	1.8	1.4	1.6	10.8
	2E+06	1.7	1.4	1.6	10.9
	1E+06	1.7	1.3	1.6	10.1
	6E+05	1.7	1.3	1.5	9.5
	2E+05	1.6	1.0	1.4	7.9

Setup components: total time, 1M dof, unstructured

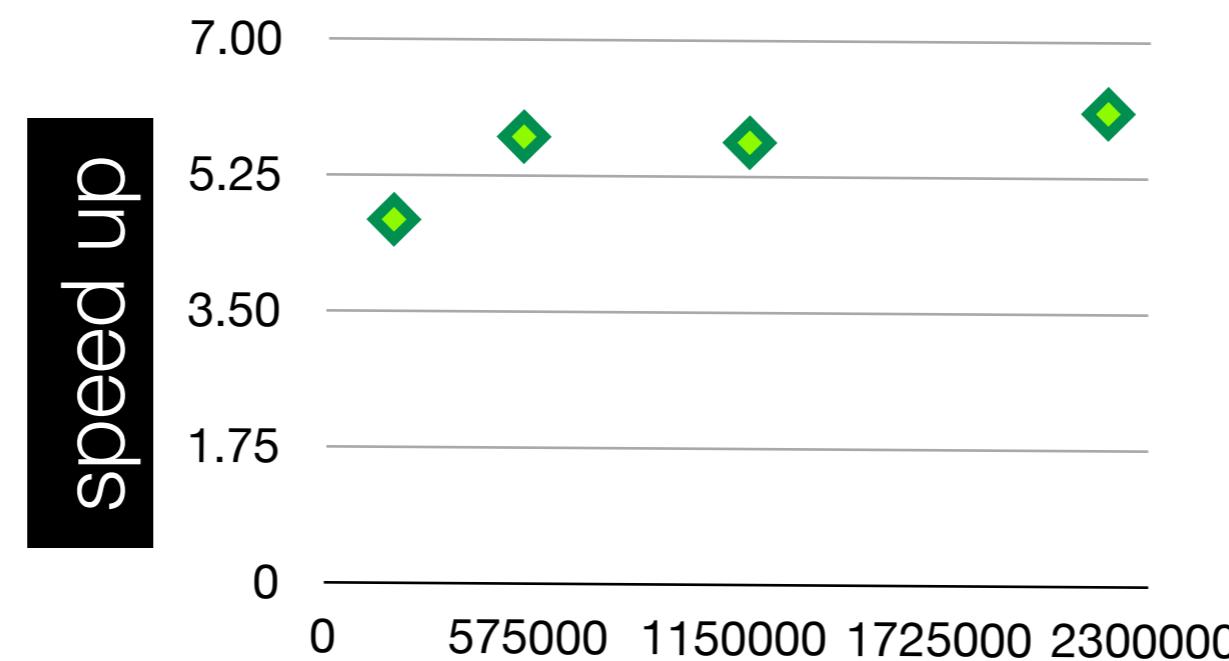
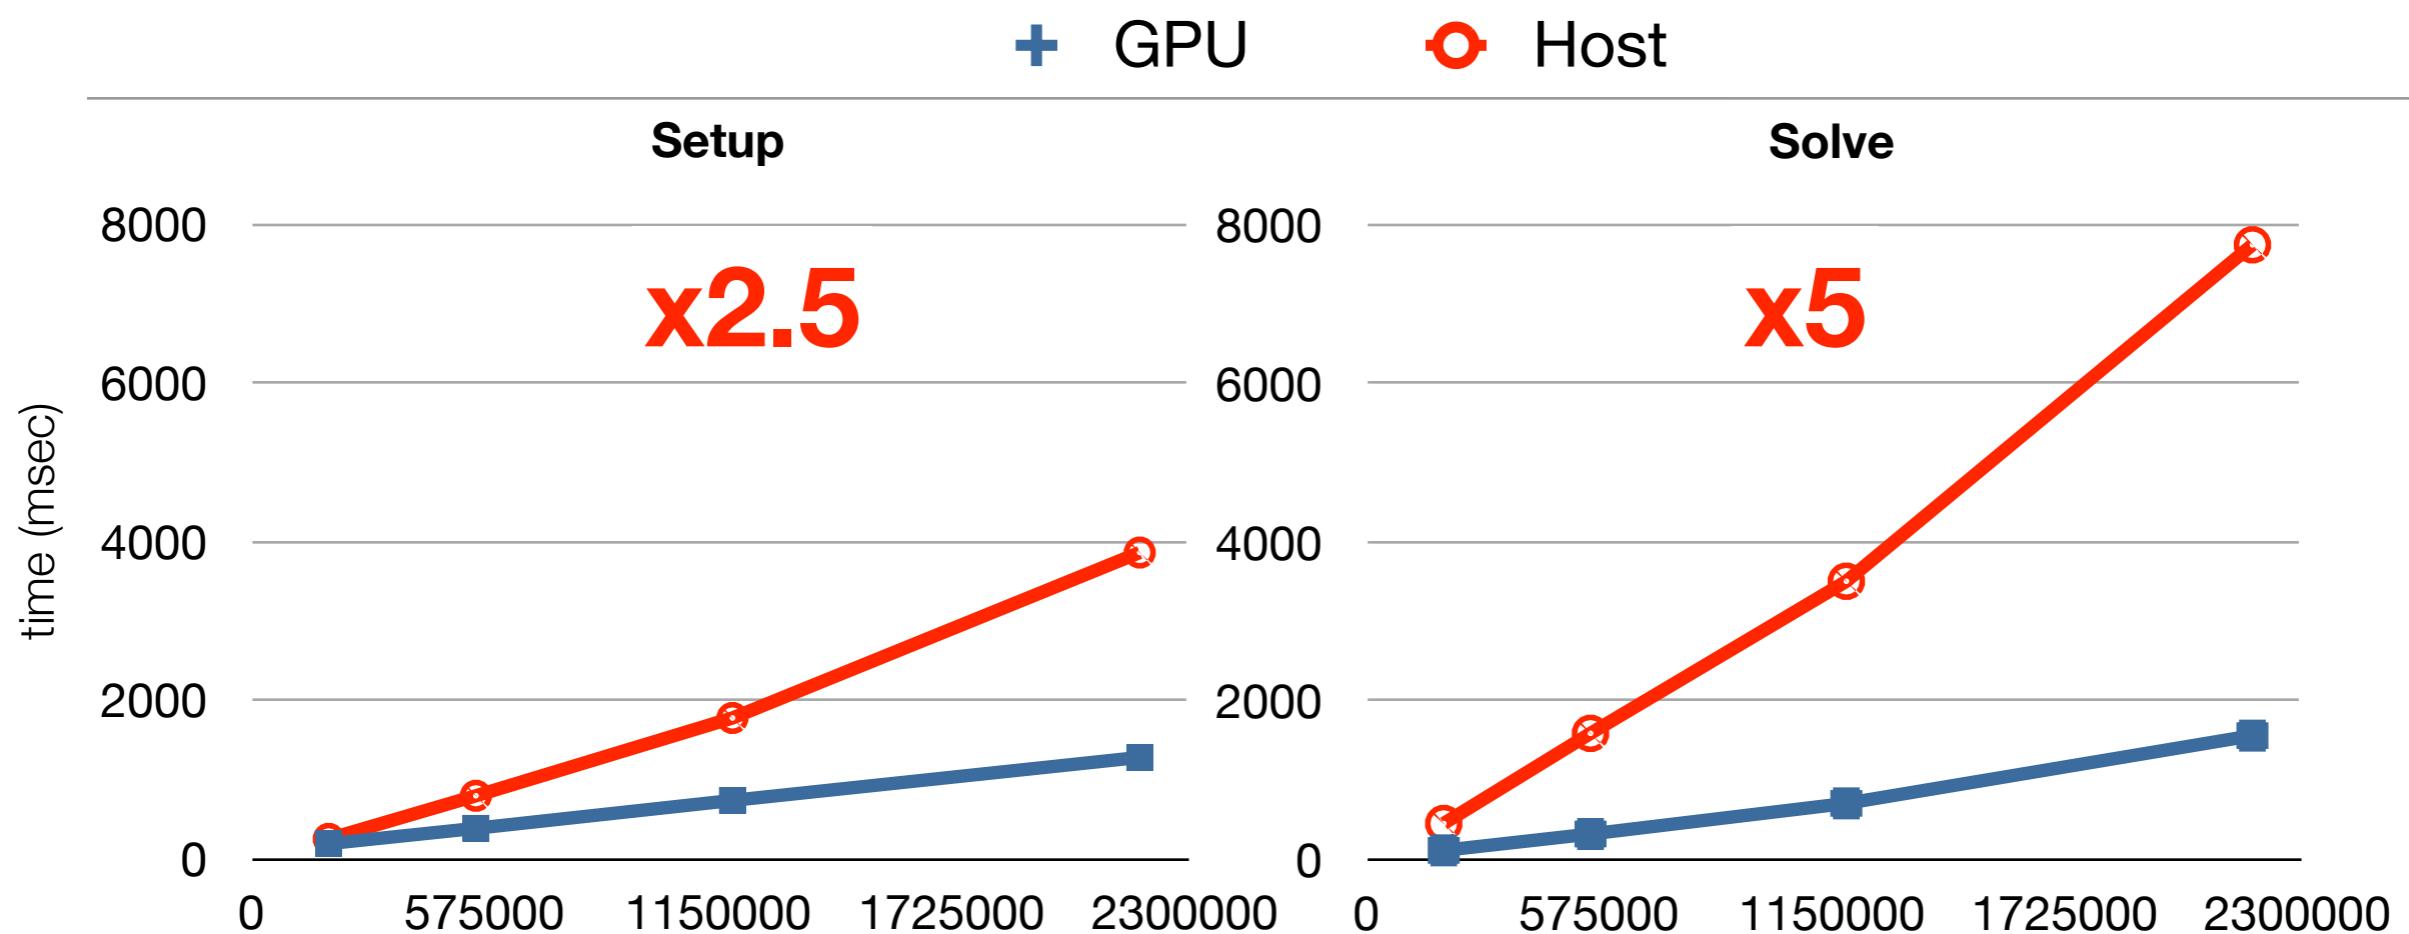


Triple products are expensive

speed up



GPU likes fine (sparse) levels



solvers and applications (discretizations)

- **goal:** a solver that can utilize natural scale in the problem
 - wave problems
 - high-order elements
 - many graph problems
- **goal:** best balance of accuracy and cost
 - LS FEM often lead to more cooperative sparse systems
 - high-order may have higher degrees of fine-grained parallelism
 - conforming elements may lead to a more natural hierarchy/null-space
- **goal:** more useable solver for applications
 - not blackbox
 - more guidance on tuning and a measure of success

solvers and theory

- **goal:** predictive theory to assess quality/efficiency *a priori*
 - sharp upper bounds on convergence
 - useful lower bounds on convergence
 - performance measures that assess the complexity
- **goal:** fully adaptive methods
 - fine-tune the solver based on history
- **goal:** ubiquitous complexity
 - globally aware decisions in the solver setup
 - architecture sensitive theory

solvers and computation

- golden age of multigrid and computation?
- **goal:** scalability
 - parallel modeling (Gahvari, Gropp, et al.) tied to development
 - additive, repetitive, or redundant cycling
- **goal:** learning, automating, tuning
 - use a scaling phase for learning
 - parameter pool
- **goal:** heterogeneous environments
 - fine-grained approach
 - flexible use of resources
- low-power?

Contributions

- NSF-DMS
- Teragrid for allocations
- Nvidia for hardware
- software development: CUSP::MG
- software development: PyAMG
- LLNL, SNL for student support

